

Bayesian Quantile Regression with Deep Echo State Network Features

Antonio De Leon¹ Raquel Prado¹ Bruno Sansó¹

¹Department of Statistics, University of California, Santa Cruz

August 15, 2026

Abstract

Conditional quantiles are central to asymmetric decision losses, tail-risk assessment, and interval forecasts, but Bayesian quantile regression for nonlinear dynamic series can be difficult when the temporal feature vector is high-dimensional. We develop the quantile deep echo state network (Q-DESN), a Bayesian quantile regression model conditional on fixed features generated by a deep echo state network. After the reservoir architecture, random seed, scaling, reducers, and washout convention are fixed, posterior uncertainty is assigned to the regression coefficients, likelihood, and shrinkage parameters. Single-level fits use asymmetric Laplace or quantile-fixed generalized asymmetric Laplace working likelihoods with ridge or regularized-horseshoe priors. Posterior computation uses Markov chain Monte Carlo where feasible and a model-specific variational Bayes approximation for lower-cost screening, initialization, and specific applications. For quantile grids, we compare independent level-wise regressions followed by monotone rearrangement with a joint quantile-vector regression that shrinks adjacent quantile-specific coefficient differences. Synthetic studies, a single-origin retrospective GloFAS streamflow case, and a retrospective PriceFM comparison identify settings where this fixed-feature Bayesian regression improves finite-grid quantile scores.

1 Introduction

Forecast evaluation often requires summaries beyond the conditional mean. Conditional means are natural under squared-error loss, but asymmetric decision losses, tail-risk assessment, and interval forecasts call for conditional quantiles. Under asymmetric piecewise-linear loss, conditional quantiles are optimal point forecasts, and pairs of quantiles define interval forecasts (Koenker and Bassett, 1978; Koenker, 2005; Gneiting, 2011). Let Y_{t+h} denote the future response at horizon h , and let $\mathcal{F}_t$ be the sigma-field generated by the information available at forecast origin t , including past responses and covariates known at the origin or specified through a scenario. With $F_{t,h}(u) = \Pr(Y_{t+h} \leq u \mid \mathcal{F}_t)$, define the level- p_0 conditional quantile by the generalized inverse

$$Q_{t,h}(p_0) = F_{t,h}^{-1}(p_0) = \inf\{u : F_{t,h}(u) \geq p_0\}, \quad p_0 \in (0, 1).$$

When the forecasting question concerns one probability level, a single fitted conditional quantile is the relevant summary. When several levels are fitted, the resulting finite grid can be rearranged into a monotone quantile curve. Such a curve is a predictive summary over the chosen grid. A full predictive distribution additionally requires a construction that couples the levels into a distribution (Gneiting and Raftery, 2007; Gneiting and Katzfuss, 2014; Barlow and Brunk, 1972; Chernozhukov et al., 2010).

For nonlinear dynamic data, quantile forecast accuracy also depends on how temporal dependence is represented. An echo state network (ESN) represents past responses and covariates through a recurrent nonlinear state equation. In the construction used here, the input and recurrent weights are generated once and then held fixed. Running this fixed recurrence through the observed history produces reservoir states, which are then used as rows of a nonlinear design matrix for regression. The ESN literature calls the final regression layer the readout (Jaeger, 2001; Lukoševičius and Jaeger, 2009; Lukoševičius, 2012). A deep echo state network (DESN) extends this idea by stacking reservoir layers, following the deep ESN construction of Gallicchio et al. (2018). Conditional on the specified architecture, random seed, scaling constants, reducers, and washout convention, the DESN therefore supplies a deterministic nonlinear feature matrix. Q-DESN uses Bayesian quantile regression with that fixed design, conditioning on the fixed recurrent weights throughout the analysis.

The quantile deep echo state network (Q-DESN) combines this fixed DESN design with a quantile working likelihood. For a chosen level p_0 , the linear predictor $\mu_t = \mathbf{x}_t^\top \boldsymbol{\beta}$ is the likelihood location and, under the parameterization used here, the fitted p_0 conditional quantile. The baseline likelihood is asymmetric Laplace (AL). The extended asymmetric Laplace likelihood, denoted exAL, is the quantile-fixed generalized asymmetric Laplace construction of Yan et al. (2025); it allows mode, skewness, and tail behavior to vary while retaining the fitted quantile level. In model labels, the prefix “ex” indicates this extended working likelihood, while unprefix QDESN and DQLM labels use AL.

The AL or exAL likelihood is used as a working likelihood in the Bayesian quantile-regression sense. It defines a posterior distribution for quantile-specific coefficients without asserting that either likelihood is the data-generating error law (Sriram et al., 2013; Yang et al., 2016; Ji et al., 2025). Posterior intervals for coefficient-derived quantiles and forecast bands computed from the fitted likelihood are therefore model-based summaries conditional on the fixed DESN design and the chosen working likelihood. Frequentist coverage of credible intervals under likelihood misspecification is a separate question from empirical coverage of the forecast bands; the simulations and retrospective applications assess the latter empirically.

High-dimensional regressions on DESN features require regularization. The default coefficient prior is ridge, which gives dense Gaussian shrinkage. As an adaptive alternative, the regularized horseshoe (RHS) introduces global-local shrinkage and a slab component for large coefficients (Carvalho et al., 2010; Piironen and Vehtari, 2017). Under the product representation used here, the local, global, and slab scale updates retain closed-form inverse-gamma forms (Nishimura and Suchard, 2023; Makalic and Schmidt, 2016). Posterior computation uses Markov chain Monte Carlo

(MCMC) as the reference method when the run is feasible. A mean-field variational Bayes (VB) approximation is used for screening, initialization, and specific applications. For non-conjugate exAL scale-asymmetry blocks, the VB implementation uses the Laplace–Delta approximation of Wang and Blei (2013); the label VB–LD refers to this model-specific approximation.

This article makes three principal contributions. First, it introduces Q–DESN as Bayesian conditional-quantile regression on a fixed DESN feature map, with AL or exAL working likelihoods and ridge or regularized-horseshoe regularization. Second, it develops two distinct strategies for quantile grids: independent level-wise regressions followed by a pre-specified monotone rearrangement, and a joint quantile-vector regression that shares information through regularized-horseshoe shrinkage on adjacent coefficient differences, borrowing from Quantile-Varying Parameter ideas (Kohns and Szendrei, 2025). This adjacent-difference prior encourages similar slopes at neighboring quantile levels; exact noncrossing still requires an explicit constraint or monotone rearrangement. Third, the article gives MCMC and model-specific VB–LD computation and evaluates the resulting procedures in simulations and two retrospective applications. The applications are interpreted under their stated covariates and forecast origins.

The empirical sections use Q–DESN in three ways. The single-quantile simulation compares QDESN and exQDESN regressions with DQLM and exDQLM baselines. The joint simulation isolates multi-quantile coefficient sharing against independent level-wise regressions under known true conditional quantile paths. The GloFAS and PriceFM examples use independent Q–DESN regressions in retrospective forecasting comparisons. Thus, evidence for the joint quantile-vector prior comes from the synthetic multi-quantile study, whereas the applications assess fixed-feature Bayesian quantile regression in two domain-specific forecasting settings.

Related Work

Bayesian quantile regression provides the likelihood-based foundation for Q–DESN. Classical quantile regression targets conditional quantiles directly (Koenker and Bassett, 1978; Koenker, 2005). Bayesian implementations often use the asymmetric Laplace likelihood because its mode aligns with check-loss minimization and because its mixture representation supports conditional posterior computation (Yu and Moyeed, 2001; Kozumi and Kobayashi, 2011). In this article the asymmetric likelihood is a working likelihood: it defines posterior uncertainty for quantile-regression parameters, while the frequentist behavior of model-based credible intervals requires separate assessment (Sriram et al., 2013; Yang et al., 2016; Ji et al., 2025). The quantile-fixed generalized asymmetric Laplace distribution of Yan et al. (2025) relaxes the AL error-shape restriction after the fitted quantile level has been fixed; this is the exAL working likelihood used below.

Dynamic quantile linear models provide Bayesian state-space baselines for time-varying conditional quantiles. The dynamic quantile linear model (DQLM) combines a linear state evolution with Bayesian quantile regression (Gonçalves et al., 2020). The extended DQLM (exDQLM) of Barata et al. (2022) adds the corresponding asymmetric-likelihood flexibility, and both paths are implemented in the CRAN package `exdqlm` (Barata et al., 2026). Q–DESN changes the representation of

temporal dependence: rather than specifying a low-dimensional linear state process, it uses a fixed nonlinear reservoir state as the regression design. The DQLM and exDQLM fits therefore serve as structured linear dynamic quantile comparators under the same held-out forecast evaluation.

Reservoir computing and ESNs provide the nonlinear dynamic features used by Q-DESN. ESNs were introduced by Jaeger (2001) and later reviewed within the broader reservoir-computing literature by Lukoševičius and Jaeger (2009); practical guidance is given by Lukoševičius (2012). In a leaky-integrator ESN, the leak rate controls how strongly the current state moves toward the new input-driven update rather than persisting from the previous state. Analyses of the echo state property caution that spectral-radius heuristics should be checked rather than treated as sufficient stability guarantees (Jaeger et al., 2007; Yildiz et al., 2012). Broader reservoir-computing surveys and time-series reviews discuss their role in forecasting and applied dynamic modeling (Zhang and Vargas, 2023; Yan et al., 2024; Cardoso et al., 2024; Sun et al., 2024). Deep ESNs, topology variants, and reservoir hyperparameter tuning extend the basic construction by stacking reservoirs, changing connectivity, or choosing reservoir parameters (Gallicchio et al., 2018; Kim and King, 2020; Maat et al., 2018; Bai et al., 2023; Viehweg et al., 2025). Q-DESN uses these reservoir constructions as fixed nonlinear features for Bayesian quantile regression.

A related reservoir-computing literature embeds ESN features in probabilistic forecasting models. Deep ESNs with uncertainty quantification, probabilistic reservoir methods for load forecasting, activation-level uncertainty propagation, and Bayesian or state-space ESN formulations show that reservoir features can support uncertainty-aware prediction (McDermott and Wikle, 2019; Guerra et al., 2023; Gandhi et al., 2018; Singh and Raman, 2025). The narrower gap addressed here is Bayesian inference for conditional quantiles with fixed DESN features, asymmetric working likelihoods, and high-dimensional coefficient regularization.

Quantile-oriented ESN methods provide the closest forecasting comparison. Lv et al. (2016) use quantile-regression ESN ensembles to construct prediction intervals for blast-furnace gas generation, and Bonas et al. (2024) use penalized quantile regression to calibrate DESN forecasts for quasi-periodic climate processes. Q-DESN is complementary to these approaches. Conditional on the generated reservoir, it specifies an exAL regression whose posterior includes coefficients, likelihood scale, asymmetry, and shrinkage parameters. The joint extension further couples quantile-specific regression coefficients during inference, so raw multi-quantile paths can be more coherent before monotone rearrangement is applied.

Shrinkage priors enter because fixed reservoir designs can be high-dimensional and correlated. Ridge shrinkage provides the default dense Gaussian regularization. The regularized horseshoe, building on the horseshoe prior and its slab-regularized extension, is used as an adaptive global-local alternative (Carvalho et al., 2010; Piironen and Vehtari, 2017). Under the product representation used here, scale updates remain closed form (Nishimura and Suchard, 2023; Makalic and Schmidt, 2016). These priors regularize the coefficients. Feature-level interpretability requires additional diagnostics.

Finally, a grid of quantile levels requires a distinction between monotone rearrangement after

fitting and joint modeling across levels. Separate Q-DESN fits can cross because each level is estimated independently. Isotonic regression can project the fitted grid onto the monotone cone (Barlow and Brunk, 1972), and rearrangement methods provide a general approach to noncrossing quantile and probability curves (Chernozhukov et al., 2010). These operations repair the displayed quantile curve after fitting. A joint posterior over quantile-indexed coefficients requires a joint model across levels. Constrained quantile-regression methods instead impose non-crossing restrictions during estimation (Wu and Liu, 2009; Bondell et al., 2010). Recent work also shows that non-crossing constraints can be viewed as fused shrinkage across quantile-specific coefficients (Szendrei et al., 2025). Bayesian joint quantile models share information across levels through process priors, score-likelihood constructions, or smoothness penalties (Yang and Tokdar, 2017; Wu and Narisetty, 2021; Wang and Cai, 2024). The Quantile-Varying Parameter construction of Kohns and Szendrei (2025) is especially relevant here because it places a state-space shrinkage prior on adjacent quantile-indexed coefficient changes. The joint Q-DESN extension below adapts that idea to fixed-feature regressions and assigns regularized-horseshoe priors to adjacent slope differences. The resulting prior encourages neighboring quantile curves to have similar slopes, which can reduce crossings. In the simulation study below, it substantially reduces raw crossings relative to independent AL regressions before monotone rearrangement; exact monotonicity still requires an explicit ordered parameterization, projection, or rearrangement.

These distinctions determine the notation used below. We first separate the objects fixed by reservoir construction from the stochastic coefficient and likelihood quantities, so that posterior statements can be read as conditional on a recorded feature map.

The remainder of the paper is organized as follows. Section 2 introduces notation, deep echo state network (DESN) dynamics, and the exAL representation. Section 3 presents the Q-DESN model and prior specification, including independent and joint multi-quantile extensions. Section 4 describes posterior computation, while Sections 5–6 describe multi-step quantile summaries, finite-grid scoring, and reservoir diagnostics. Section 7 presents the simulation design, and Sections 8–9 give streamflow and electricity-price forecasting applications. Section 10 summarizes limitations and future work. The supplement gives full conditional derivations, variational and ELBO details, monotone rearrangement details, and supplementary simulation and application tables.

2 Notation and Preliminaries

The notation separates deterministic reservoir quantities from stochastic coefficient and likelihood quantities. Matrices and vectors are bold, as in $\mathbf{W}$ and $\mathbf{x}$; scalar observations, latent states, and parameters use standard lowercase notation unless they are named constants. The symbols IG, GIG, and $\mathcal{N}^+$ denote inverse-gamma, generalized inverse Gaussian, and positive-truncated Normal distributions. For a square matrix $\mathbf{A}$, write $\rho(\mathbf{A}) = \max_i |\lambda_i(\mathbf{A})|$, where $\{\lambda_i(\mathbf{A})\}$ are the eigenvalues. Thus $\rho(\cdot)$ denotes spectral radius, while $\rho_d \in (0, 1)$ denotes the target spectral radius for layer d . Indices use $d = 1, \dots, D$ and $t = 1, \dots, T$. When conditioning information is shown explicitly,

$Q_p(y \mid \cdot)$ denotes the conditional p -quantile of the scalar response. The identity matrix, zero vector, and one vector of dimension m are denoted by $\mathbf{I}_m$, $\mathbf{0}_m$, and $\mathbf{1}_m$, respectively.

2.1 Deep Echo State Network (DESN)

The DESN is a deterministic feature map once the architecture, random seed, scaling constants, reducers, and washout convention are fixed. Conditional on these choices, posterior inference is restricted to the quantile regression model. The following notation states the dimensions before the recursion so that each layer-specific matrix has a defined statistical role.

Notation and dimensions. The input vector is $\mathbf{u}_t = (y_{t-1}, \dots, y_{t-m}, \mathbf{z}_t^\top)^\top \in \mathbb{R}^{m+q}$, with m response lags and q exogenous variables. For layer d , $\mathbf{W}_d \in \mathbb{R}^{n_d \times n_d}$ is recurrent, $\mathbf{W}_1^{\text{in}} \in \mathbb{R}^{n_1 \times (m+q)}$ maps the original input to the first layer, and $\mathbf{W}_d^{\text{in}} \in \mathbb{R}^{n_d \times \tilde{n}_{d-1}}$ maps the reduced state from layer $d-1$ to layer d for $d \geq 2$. The reducer $\mathbf{Q}_{d-1} \in \mathbb{R}^{\tilde{n}_{d-1} \times n_{d-1}}$ maps $\mathbf{h}_{t,d-1} \in \mathbb{R}^{n_{d-1}}$ to $\tilde{\mathbf{h}}_{t,d-1} \in \mathbb{R}^{\tilde{n}_{d-1}}$. The regression dimension is $r = n_D + \sum_{d=1}^{D-1} \tilde{n}_d + (m+q)$, and $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_r)^\top \in \mathbb{R}^{r+1}$.

Random sparse weights and spectral scaling. In contrast to fully trained recurrent networks, reservoir and input weights are drawn once and kept fixed throughout inference. Let $\circ$ denote the Hadamard product. Each recurrent matrix is $\mathbf{W}_d = \Phi_d^{(W)} \circ Z_d^{(W)}$, and each input matrix is generated analogously as $\mathbf{W}_d^{\text{in}} = \Phi_d^{(\text{in})} \circ Z_d^{(\text{in})}$. Mask entries are Bernoulli with probabilities $\pi_d^{(W)}$ and $\pi_d^{(\text{in})}$; base entries are independent draws from zero-mean finite-variance distributions p_W and p_{in} , such as $\mathcal{N}(0, \sigma^2)$ or $\text{Unif}[-a, a]$. Masks and base draws are mutually independent and independent across entries and layers. The recurrent matrices are then scaled to the target spectral radii,

$$\text{Spectral-radius scaling:} \quad \bar{\mathbf{W}}_d = \frac{\rho_d}{\rho(\mathbf{W}_d)} \mathbf{W}_d, \quad \rho_d \in (0, 1).$$

If $\rho(\mathbf{W}_d) = 0$, the matrix is resampled. Input matrices $\mathbf{W}_d^{\text{in}}$ are not spectrally scaled.

DESN recursion and regression features. With layer-specific leak rates $\alpha_d \in [0, 1]$, elementwise nonlinearity $f(\cdot)$, and lower-layer transformation $k(\cdot)$, the fixed feature map is

$$\begin{aligned} \text{Layer 1 innovation:} \quad & \boldsymbol{\omega}_{t,1} = f\left(\bar{\mathbf{W}}_1 \mathbf{h}_{t-1,1} + \mathbf{W}_1^{\text{in}} \mathbf{u}_t\right), \\ \text{Layer 1 update:} \quad & \mathbf{h}_{t,1} = (1 - \alpha_1) \mathbf{h}_{t-1,1} + \alpha_1 \boldsymbol{\omega}_{t,1}, \\ \text{Reduced bridge state:} \quad & \tilde{\mathbf{h}}_{t,d-1} = \mathbf{Q}_{d-1} \mathbf{h}_{t,d-1}, \quad d = 2, \dots, D, \\ \text{Layer } d \text{ innovation:} \quad & \boldsymbol{\omega}_{t,d} = f\left(\bar{\mathbf{W}}_d \mathbf{h}_{t-1,d} + \mathbf{W}_d^{\text{in}} \tilde{\mathbf{h}}_{t,d-1}\right), \quad d = 2, \dots, D, \\ \text{Layer } d \text{ update:} \quad & \mathbf{h}_{t,d} = (1 - \alpha_d) \mathbf{h}_{t-1,d} + \alpha_d \boldsymbol{\omega}_{t,d}, \quad d = 2, \dots, D, \\ \text{Feature stack:} \quad & \tilde{\mathbf{x}}_t = [\mathbf{h}_{t,D}^\top, k(\tilde{\mathbf{h}}_{t,1})^\top, \dots, k(\tilde{\mathbf{h}}_{t,D-1})^\top, \mathbf{u}_t^\top]^\top, \\ \text{Intercept and linear predictor:} \quad & \mathbf{x}_t = (1, \tilde{\mathbf{x}}_t^\top)^\top, \quad \mu_t = \mathbf{x}_t^\top \boldsymbol{\beta}. \end{aligned} \tag{2.1}$$

Here $k(\cdot)$ is an elementwise activation applied to reduced lower-layer states, commonly $k = \text{id}$ or $k = f$.

The input vector $\mathbf{u}_t$ bundles response lags and exogenous covariates; the intercept is appended only in $\mathbf{x}_t$. The leaky update is a convex combination of the previous state and the current nonlinear update: small α_d retains more memory, while large α_d reacts more quickly to the current input. The leak rate α_d and target radius ρ_d therefore control the empirical time scale and persistence of the deterministic state recursion; they are tuning choices rather than sufficient guarantees of the echo state property. Reducers $\mathbf{Q}_{d-1}$ may be fixed random projections, optionally orthonormalized, or precomputed linear reducers such as PCA from an initial preprocessing pass.

The feature stack concatenates the top-layer state, transformed reduced lower-layer states, and $\mathbf{u}_t$, yielding $\tilde{\mathbf{x}}_t \in \mathbb{R}^r$. Appending an intercept gives $\mathbf{x}_t \in \mathbb{R}^{r+1}$, and the coefficient vector maps this feature vector to the scalar location μ_t . States are initialized at $\mathbf{h}_{0,d} = \mathbf{0}_{n_d}$, or from $\mathcal{N}(\mathbf{0}_{n_d}, \sigma_{h,d}^2 \mathbf{I}_{n_d})$. A teacher-forcing washout $t = 1, \dots, T_0$ is run before features enter the regression; this discards early transient states that depend strongly on the arbitrary initialization. Washout length and scaling choices can affect short-horizon behavior. All masks and base matrices are fixed for the entire analysis and are not resampled during posterior inference.

2.2 Extended Asymmetric Laplace Working Likelihood

The quantile-fixed generalized asymmetric Laplace distribution of Yan et al. (2025) extends the asymmetric Laplace (AL) working likelihood while keeping the target quantile fixed. We denote this likelihood by exAL, for extended asymmetric Laplace. The standard AL likelihood is recovered as a special case, but its skewness is fixed once p_0 is chosen. The exAL asymmetry parameter allows mode, skewness, and tail behavior to vary within a quantile-fixed likelihood. At fixed level p_0 , μ is the p_0 th quantile, $\sigma > 0$ is a scale parameter, and γ is an asymmetry parameter with bounded support (L, U) . The endpoints L and U solve $g(\gamma) = 1 - p_0$ and $g(\gamma) = p_0$, respectively, where $g(\gamma) = 2\Phi(-|\gamma|)\exp(\gamma^2/2)$, and Φ is the standard Normal CDF.

Posterior computation uses the stochastic representation of the exAL distribution. If

$$y \sim \text{exAL}_{p_0}(\mu, \sigma, \gamma),$$

then there exist independent $\epsilon \sim \mathcal{N}(0, 1)$, $z \sim \text{Exp}(1)$, and $s \sim \mathcal{N}^+(0, 1)$ such that

$$y = \mu + \lambda(\gamma)\sigma s + A(\gamma)\sigma z + \{B(\gamma)\sigma^2 z\}^{1/2}\epsilon,$$

where $A(\gamma) = \{1 - 2p_\gamma\}/\{p_\gamma(1 - p_\gamma)\}$, $B(\gamma) = 2/\{p_\gamma(1 - p_\gamma)\}$, $C(\gamma) = \{\mathbf{1}\{\gamma > 0\} - p_\gamma\}^{-1}$, $p_\gamma = \mathbf{1}\{\gamma < 0\} + \{p_0 - \mathbf{1}\{\gamma < 0\}\}/g(\gamma)$, and $\lambda(\gamma) = C(\gamma)|\gamma|$. Setting $v = \sigma z$ gives $v \mid \sigma \sim \text{Exp}(\text{rate} = 1/\sigma)$ and the conditional Gaussian form used for Q-DESN posterior computation:

$$y \mid \mu, \sigma, \gamma, s, v \sim \mathcal{N}(\mu + \lambda(\gamma)\sigma s + A(\gamma)v, B(\gamma)\sigma v).$$

All quantities $g(\gamma)$, p_γ , $A(\gamma)$, $B(\gamma)$, $C(\gamma)$, and $\lambda(\gamma)$ depend on the target quantile level p_0 ; this dependence is suppressed for readability. Once (p_0, γ) is fixed, these are deterministic constants in

the Gaussian augmentation.

3 Model Specification

Fix a quantile level $p_0 \in (0, 1)$. After reservoir construction and washout, $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_T)^\top$ is treated as fixed, and all posterior uncertainty is conditional on this design. The statistical model is a Bayesian quantile regression for $\mathbf{y}$ given $\mathbf{X}$. The observation model and priors define the posterior distribution; the Gaussian augmentation below is a computational representation of the asymmetric working likelihood. Dependence on p_0 is suppressed when no ambiguity arises.

3.1 Quantile Deep Echo State Network (Q-DESN)

The Q-DESN is obtained by inserting the fixed DESN feature map into the exAL working likelihood of Section 2.2. Let $\tilde{\mathbf{x}}_t$ be the feature vector in Eq. (2.1), and set $\mathbf{x}_t = (1, \tilde{\mathbf{x}}_t^\top)^\top$. Conditional on the fixed design matrix, the marginal working likelihood is

$$y_t \mid \mathbf{x}_t, \boldsymbol{\beta}, \sigma, \gamma \sim \text{exAL}_{p_0}(\mathbf{x}_t^\top \boldsymbol{\beta}, \sigma, \gamma), \quad t = 1, \dots, T.$$

Thus the dynamic linear predictor $\mathbf{x}_t^\top \boldsymbol{\beta}$ is the model's conditional p_0 -quantile, while σ and γ control the scale and shape of the working error distribution. For posterior computation, the same exAL mixture representation is applied at each time point:

$$\begin{aligned} \text{Quantile linear predictor :} \quad & \mu_t = \mathbf{x}_t^\top \boldsymbol{\beta}, \quad \boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_r)^\top \in \mathbb{R}^{r+1}, \\ \text{Gaussian conditional :} \quad & y_t \mid \boldsymbol{\beta}, \sigma, \gamma, s_t, v_t, \mathbf{x}_t \sim \mathcal{N}(\mu_t + \lambda(\gamma)\sigma s_t + A(\gamma)v_t, B(\gamma)\sigma v_t), \\ \text{Positive-shift latent :} \quad & s_t \sim \mathcal{N}^+(0, 1), \\ \text{Scale-mixture latent :} \quad & v_t \mid \sigma \sim \text{Exp}(\text{rate} = 1/\sigma). \end{aligned} \tag{3.1}$$

The display separates the statistical quantile model from the computational augmentation. The first line defines the quantile path; the second gives the conditional Gaussian observation equation; and the last two lines introduce auxiliary variables. The constants $A(\gamma)$, $B(\gamma)$, and $\lambda(\gamma)$ are inherited from Section 2.2. The DESN feature map is fixed before posterior inference, while $\boldsymbol{\beta}$, $\sigma > 0$, and $\gamma \in (L, U)$ are posterior unknowns. Integrating out (s_t, v_t) recovers the marginal exAL likelihood. These variables expose conditionally Gaussian updates while preserving the quantile target. Complete-data joint distributions for the AL and exAL likelihoods with ridge and RHS priors are given in Section S4 of the supplement.

3.2 Single-Level Prior Specification

The intercept is modeled separately from the slope coefficients: $\beta_0 \sim \mathcal{N}(b_0, V_0)$, with $b_0 \in \mathbb{R}$ and $V_0 > 0$. This lets shrinkage act on the high-dimensional DESN features without forcing the overall

fitted quantile level toward zero.

Ridge prior. For the non-intercept coefficients $\beta_{\text{slope}} = (\beta_1, \dots, \beta_r)^\top$, the default ridge prior is

$$\beta_{\text{slope}} \mid \kappa_\beta^2 \sim \mathcal{N}(\mathbf{0}_r, \kappa_\beta^2 \mathbf{I}_r), \quad (3.2)$$

where $\kappa_\beta^2 > 0$ is fixed or assigned a weakly informative scale prior. This Gaussian prior gives dense baseline regularization.

Adaptive shrinkage. The regularized horseshoe (RHS) prior is the adaptive shrinkage alternative. It starts from the ordinary horseshoe (Carvalho et al., 2010), adds slab regularization (Piironen and Vehtari, 2017), and uses the product representation of Nishimura and Suchard (2023) to retain closed-form scale updates. For non-intercept coefficients $j = 1, \dots, r$, the ordinary horseshoe is

$$\beta_j \mid \tau, \lambda_j \sim \mathcal{N}(0, \tau^2 \lambda_j^2), \quad \lambda_j \sim C^+(0, 1), \quad \tau \sim C^+(0, \tau_0), \quad (3.3)$$

where τ is the global scale and λ_j is coefficient-specific. This prior concentrates mass near zero while allowing some coefficients to escape strong shrinkage. The regularized version introduces a slab scale $\zeta > 0$ to control the largest prior variances and sets

$$V_j(\lambda_j, \tau, \zeta) = \left(\zeta^{-2} + \tau^{-2} \lambda_j^{-2} \right)^{-1} = \frac{\zeta^2 \tau^2 \lambda_j^2}{\zeta^2 + \tau^2 \lambda_j^2}, \quad j = 1, \dots, r, \quad (3.4)$$

with conditional prior

$$\beta_j \mid \lambda_j, \tau, \zeta \sim \mathcal{N}(0, V_j), \quad j = 1, \dots, r. \quad (3.5)$$

This is the RHS coefficient prior used in the adaptive-shrinkage Q-DESN fits. Equation (3.4) gives the interpretation: if $\tau^2 \lambda_j^2 \ll \zeta^2$, then $V_j \approx \tau^2 \lambda_j^2$, recovering ordinary horseshoe behavior; if $\tau^2 \lambda_j^2 \gg \zeta^2$, then $V_j \approx \zeta^2$, so the slab caps the effective prior variance.

The regularized horseshoe can be written directly as in Piironen and Vehtari (2017). For Gibbs and coordinate-ascent variational updates, however, the direct global-local block includes an extra $\{1 + \tau^2 \lambda_j^2 / \zeta^2\}^{1/2}$ factor, so the ordinary horseshoe updates for (τ, λ_j) are not preserved. We therefore use the equivalent product representation of Nishimura and Suchard (2023), which keeps the same conditional coefficient prior while restoring closed-form scale updates:

$$p(\beta_j, \lambda_j \mid \tau, \zeta) \propto \exp\left(-\frac{\beta_j^2}{2\tau^2 \lambda_j^2}\right) \exp\left(-\frac{\beta_j^2}{2\zeta^2}\right) \lambda_j^{-1} (1 + \lambda_j^2)^{-1}, \quad (3.6)$$

where proportionality is with respect to (β_j, λ_j) . The (τ, ζ) -dependent Gaussian normalizing factors are retained in the full-conditionals derivations so that the inverse-gamma shape parameters are recovered correctly. Thus the RHS option gives adaptive global-local shrinkage while retaining inverse-gamma updates for the local and global scales.

Closed-form scale updates use the inverse-gamma representation of half-Cauchy priors (Makalic

and Schmidt, 2016):

$$\lambda_j^2 \mid \nu_j \sim \text{IG}\left(\frac{1}{2}, \frac{1}{\nu_j}\right), \quad \nu_j \sim \text{IG}\left(\frac{1}{2}, 1\right), \quad j = 1, \dots, r, \quad (3.7)$$

$$\tau^2 \mid \xi \sim \text{IG}\left(\frac{1}{2}, \frac{1}{\xi}\right), \quad \xi \sim \text{IG}\left(\frac{1}{2}, \frac{1}{\tau_0^2}\right). \quad (3.8)$$

The slab scale is either fixed or assigned

$$\zeta^2 \sim \text{IG}(a_\zeta, b_\zeta). \quad (3.9)$$

Under (3.6), this prior remains compatible with closed-form updating (Nishimura and Suchard, 2023).

The likelihood parameters are assigned $\sigma \sim \text{IG}(a_\sigma, b_\sigma)$ and $\gamma \sim \pi_\gamma(\gamma) \mathbf{1}\{L < \gamma < U\}$, with (L, U) determined by the quantile-fixing construction as described in Section 2.2. For RHS fits, collect $\Theta = (\beta, \sigma, \gamma, \lambda^2, \tau^2, \nu, \xi, \zeta^2)$, where $\lambda^2 = (\lambda_1^2, \dots, \lambda_r^2)^\top$ and $\nu = (\nu_1, \dots, \nu_r)^\top$. The entry ζ^2 is omitted when the slab scale is fixed. Ridge fits replace the global-local shrinkage block by κ_β^2 when that scale is assigned a prior.

3.3 Joint Quantile-Vector Regression with DESN Features

The single-level model can be fit independently over a grid of quantile levels, with separate regression coefficients at each level. When several conditional quantiles are estimated together, let $0 < p_1 < \dots < p_Q < 1$ be the fixed quantile grid and let $\mathbf{Z}$ denote the $T \times r$ non-intercept part of the fixed design matrix. For Q-DESN, $\mathbf{Z}$ is obtained from the DESN feature matrix after removing the intercept column. The joint quantile-vector regression separates intercepts from high-dimensional slopes:

$$Q_{p_q}(y_t \mid \mathcal{F}_t) = \alpha_q + \tilde{\mathbf{x}}_t^\top \beta_q, \quad q = 1, \dots, Q, \quad t = 1, \dots, T, \quad (3.10)$$

where $\alpha_q \in \mathbb{R}$ is the quantile-specific intercept and $\beta_q \in \mathbb{R}^r$ is the slope vector at level p_q . By default, the high-dimensional RHS path is placed on the slope vectors; intercepts are either weakly regularized independently or represented through ordered gaps when a stronger monotonicity convention is desired.

This construction generalizes the single-level RHS Q-DESN by replacing the single intercept-slope coefficient vector $(\beta_0, \beta_{\text{slope}})$ with $\{(\alpha_q, \beta_q)\}_{q=1}^Q$. For $Q = 1$, the individual RHS Q-DESN is recovered by taking α_1 as the single-level intercept and identifying β_1 with the single-level non-intercept slope vector, with no learned slope baseline. Equivalently, set $\beta_{\text{base}} \equiv \mathbf{0}_r$ and place the RHS prior directly on $\Delta_1 = \beta_1$.

Composite working likelihood. For each q , the working likelihood uses the exAL family at the corresponding target level,

$$y_t \mid \alpha_q, \beta_q, \sigma_q, \gamma_q \sim \text{exAL}_{p_q}\left(\alpha_q + \tilde{\mathbf{x}}_t^\top \beta_q, \sigma_q, \gamma_q\right). \quad (3.11)$$

The Q likelihood contributions form a composite working likelihood over the same response series. The joint analyses in this article use this untempered composite posterior. It updates the quantile-indexed coefficient path. Future-response simulation requires the separate forecast construction in Section 5.1.

Shrinkage across adjacent quantile levels. For a multi-quantile grid, the joint prior adapts the Quantile-Varying Parameter idea that slopes at neighboring quantile levels should be similar unless the data support tail-specific differences (Kohns and Szendrei, 2025). Let $\beta_{\text{base}} \in \mathbb{R}^r$ be a baseline slope vector and define adjacent coefficient differences

$$\Delta_{1,j} = \beta_{1,j} - \beta_{\text{base},j}, \quad \Delta_{q,j} = \beta_{q,j} - \beta_{q-1,j}, \quad q = 2, \dots, Q.$$

We assign an ordinary RHS prior to the baseline slope and RHS priors with quantile-gap-specific global scales to the adjacent coefficient differences:

$$\beta_{\text{base},j} \mid \lambda_j^0, \tau^0, \zeta_0 \sim \mathcal{N}(0, V_j^0), \quad V_j^0 = \frac{\zeta_0^2 (\tau^0)^2 (\lambda_j^0)^2}{\zeta_0^2 + (\tau^0)^2 (\lambda_j^0)^2}, \quad (3.12)$$

$$\Delta_{q,j} \mid \lambda_{q,j}^\Delta, \tau_q^\Delta, \zeta_\Delta \sim \mathcal{N}(0, V_{q,j}^\Delta), \quad V_{q,j}^\Delta = \frac{\zeta_\Delta^2 (\tau_q^\Delta)^2 (\lambda_{q,j}^\Delta)^2}{\zeta_\Delta^2 + (\tau_q^\Delta)^2 (\lambda_{q,j}^\Delta)^2}. \quad (3.13)$$

This construction separates globally useful feature effects β_{base} from effects that vary across adjacent quantile levels. It is related to interquantile shrinkage and composite-quantile regularization in frequentist quantile regression (Jiang et al., 2014; Zou and Yuan, 2008), but the shrinkage is placed on the Bayesian quantile-coefficient path conditional on the fixed feature design.

The next display rewrites the Q augmented likelihood contributions as one stacked Gaussian regression conditional on the exAL latent variables and the quantile-specific intercepts. This computational representation is algebraically equivalent to the joint quantile model above. Let $\beta_{\text{st}} = (\beta_1^\top, \dots, \beta_Q^\top)^\top$, $\mathbf{Z}_{\text{st}} = \mathbf{I}_Q \otimes \mathbf{Z}$, and $\alpha_{\text{st}} = (\alpha_1 \mathbf{1}_T^\top, \dots, \alpha_Q \mathbf{1}_T^\top)^\top$. Under the exAL augmentation, define

$$d_{q,t} = \lambda(\gamma_q) \sigma_q s_{q,t} + A(\gamma_q) v_{q,t}, \quad \omega_{q,t} = B(\gamma_q) \sigma_q v_{q,t},$$

and stack these quantities into $\mathbf{d}$ and diagonal $\mathbf{\Omega} = \text{diag}(\omega_{q,t})$. The conditional Gaussian likelihood is

$$\mathbf{y}_{\text{st}} \mid \beta_{\text{st}}, \alpha, \mathbf{d}, \mathbf{\Omega} \sim \mathcal{N}(\alpha_{\text{st}} + \mathbf{Z}_{\text{st}} \beta_{\text{st}} + \mathbf{d}, \mathbf{\Omega}), \quad \mathbf{y}_{\text{st}} = \mathbf{1}_Q \otimes \mathbf{y}.$$

Let $\mathbf{H}$ be the block first-difference matrix such that

$$\mathbf{H} \beta_{\text{st}} = (\beta_1^\top, (\beta_2 - \beta_1)^\top, \dots, (\beta_Q - \beta_{Q-1})^\top)^\top.$$

With $\tilde{\beta}_{\text{base}} = (\beta_{\text{base}}^\top, \mathbf{0}_r^\top, \dots, \mathbf{0}_r^\top)^\top$ and diagonal $\mathbf{D}_\Delta$ containing the RHS adjacent-difference variances $V_{q,j}^\Delta$, the conditional prior is proportional to

$$\exp \left[-\frac{1}{2} (\mathbf{H} \beta_{\text{st}} - \tilde{\beta}_{\text{base}})^\top \mathbf{D}_\Delta^{-1} (\mathbf{H} \beta_{\text{st}} - \tilde{\beta}_{\text{base}}) \right].$$

Consequently, the Gaussian full conditional for the stacked slope path has precision

$$\mathbf{K}_\beta = \mathbf{Z}_{\text{st}}^\top \boldsymbol{\Omega}^{-1} \mathbf{Z}_{\text{st}} + \mathbf{H}^\top \mathbf{D}_\Delta^{-1} \mathbf{H}, \quad (3.14)$$

and mean

$$\mathbf{m}_\beta = \mathbf{K}_\beta^{-1} \left[\mathbf{Z}_{\text{st}}^\top \boldsymbol{\Omega}^{-1} (\mathbf{y}_{\text{st}} - \boldsymbol{\alpha}_{\text{st}} - \mathbf{d}) + \mathbf{H}^\top \mathbf{D}_\Delta^{-1} \tilde{\boldsymbol{\beta}}_{\text{base}} \right]. \quad (3.15)$$

The sparse, block-banded prior contribution in (3.14) is the central computational object for the joint extension.

Soft noncrossing. The non-crossing interpretation is local in adjacent quantile gaps. For neighboring levels,

$$Q_{p_q}(y_t \mid \mathcal{F}_t) - Q_{p_{q-1}}(y_t \mid \mathcal{F}_t) = (\alpha_q - \alpha_{q-1}) + \tilde{\mathbf{x}}_t^\top (\boldsymbol{\beta}_q - \boldsymbol{\beta}_{q-1}).$$

On a bounded feature domain, for example after an explicit scaling or clipping condition that puts each component of $\tilde{\mathbf{x}}_t$ in $[-1, 1]$, the inequality $\alpha_q - \alpha_{q-1} \geq \|\boldsymbol{\beta}_q - \boldsymbol{\beta}_{q-1}\|_1$ is sufficient for non-crossing. The adjacent-difference RHS prior shrinks $\boldsymbol{\beta}_q - \boldsymbol{\beta}_{q-1}$ toward zero, reducing the size of the feature-dependent term and concentrating the posterior near coherent quantile paths. Exact monotonicity still requires an explicit ordered-intercept parameterization, posterior rejection rule, isotonic projection, or monotone rearrangement. The synthetic joint study below reports raw crossings separately from the pre-specified monotone rearrangement used for scored quantile grids.

4 Posterior Inference

4.1 Single-Level Augmented Posterior

For one probability level p_0 , the Q-DESN posterior conditions on the fixed DESN design and updates only the regression coefficients, working-likelihood, and shrinkage parameters. Let $\mathbf{Z}$ denote the non-intercept design, α the intercept, and $\boldsymbol{\beta}_1$ the non-intercept slope vector. Let $\boldsymbol{\Theta}_\Delta$ collect the local, global, auxiliary, and slab scales for the RHS prior when the adaptive prior is used. For the exAL working likelihood, let $\{v_t, s_t : t = 1, \dots, T\}$ be the scale-mixture and positive-shift latent variables; for AL, the s_t and γ blocks are omitted. With $\mathbf{d}_1$, $\boldsymbol{\Omega}_1$, and ω_t collecting the corresponding complete-data shifts and variances from the augmented likelihood, the single-level complete-data posterior can be written, up to constants not depending on unknowns, as

$$\begin{aligned} p(\alpha, \boldsymbol{\beta}_1, \sigma, \gamma, \{v_t, s_t\}_{t=1}^T, \boldsymbol{\Theta}_\Delta \mid \mathbf{y}, \mathbf{Z}) &\propto \exp \left[-\frac{1}{2} (\mathbf{y} - \alpha \mathbf{1}_T - \mathbf{Z} \boldsymbol{\beta}_1 - \mathbf{d}_1)^\top \boldsymbol{\Omega}_1^{-1} (\mathbf{y} - \alpha \mathbf{1}_T - \mathbf{Z} \boldsymbol{\beta}_1 - \mathbf{d}_1) \right] \\ &\times \prod_{t=1}^T \left\{ (\omega_t)^{-1/2} p(v_t \mid \sigma) p(s_t) \right\} p(\sigma) \pi_\gamma(\gamma) \mathbf{1}\{L < \gamma < U\} p(\alpha) p_\Delta^{(1)}(\boldsymbol{\beta}_1 \mid \mathbf{D}_\Delta) p(\boldsymbol{\Theta}_\Delta). \end{aligned} \quad (4.1)$$

Here

$$p_\Delta^{(1)}(\boldsymbol{\beta}_1 \mid \mathbf{D}_\Delta) \propto \exp \left(-\frac{1}{2} \boldsymbol{\beta}_1^\top \mathbf{D}_\Delta^{-1} \boldsymbol{\beta}_1 \right)$$

is the ordinary RHS slope prior, with $\Delta_1 = \beta_1$. The ridge case replaces this factor by the Gaussian ridge prior. This display gives the posterior distribution for the individual Q-DESN regression; the joint quantile-vector construction below reuses the same augmented Gaussian structure after stacking levels and replacing the direct slope prior by a regularized-horseshoe prior on quantile-indexed slope increments.

4.2 Joint Quantile-Vector Extension

For a grid $p_1 < \dots < p_Q$, let Θ_Δ collect the RHS scales for adjacent coefficient differences and let Θ_0 collect the baseline-slope RHS scales. Let Θ_Q collect α , β_{st} , the likelihood parameters, Θ_Δ , and, for $Q > 1$, $(\beta_{\text{base}}, \Theta_0)$. Let $\mathbf{a}_Q = \{v_{q,t}, s_{q,t} : q = 1, \dots, Q; t = 1, \dots, T\}$ collect the exAL auxiliary variables. With the stacked quantities $\mathbf{y}_{\text{st}}$, α_{st} , $\mathbf{Z}_{\text{st}}$, $\mathbf{d}$, and Ω defined in Section 3.3, the joint complete-data posterior is

$$\begin{aligned} p(\Theta_Q, \mathbf{a}_Q \mid \mathbf{y}, \mathbf{Z}) &\propto \exp \left[-\frac{1}{2} (\mathbf{y}_{\text{st}} - \alpha_{\text{st}} - \mathbf{Z}_{\text{st}} \beta_{\text{st}} - \mathbf{d})^\top \Omega^{-1} (\mathbf{y}_{\text{st}} - \alpha_{\text{st}} - \mathbf{Z}_{\text{st}} \beta_{\text{st}} - \mathbf{d}) \right] \\ &\times \prod_{q=1}^Q \prod_{t=1}^T \left\{ (\omega_{q,t})^{-1/2} p(v_{q,t} \mid \sigma_q) p(s_{q,t}) \right\} \prod_{q=1}^Q p(\sigma_q) \pi_{\gamma_q}(\gamma_q) \mathbf{1}\{L_q < \gamma_q < U_q\} \\ &\times p(\alpha) p_\Delta^{(Q)}(\beta_{\text{st}} \mid \beta_{\text{base}}, \mathbf{D}_\Delta) p(\Theta_\Delta) \mathcal{B}_Q, \end{aligned} \quad (4.2)$$

where

$$\mathcal{B}_Q = \begin{cases} p(\beta_{\text{base}} \mid \Theta_0) p(\Theta_0), & Q > 1, \\ 1, & Q = 1. \end{cases}$$

The Gaussian prior factor $p_\Delta^{(Q)}$ is induced by the first-difference matrix $\mathbf{H}$ and difference covariance $\mathbf{D}_\Delta$ in (3.14)–(3.15). For $Q > 1$, it is the adjacent-difference prior centered on the baseline slope, so the β_{base} argument is active. The AL special case omits $s_{q,t}$, γ_q , and the positive-truncated Normal factors. Setting $Q = 1$, taking $\alpha_1 = \alpha$, $\mathbf{y}_{\text{st}} = \mathbf{y}$, $\mathbf{Z}_{\text{st}} = \mathbf{Z}$, $\alpha_{\text{st}} = \alpha \mathbf{1}_T$, $\beta_{\text{st}} = \beta_1$, and dropping $\mathcal{B}_Q$ recovers the single-level posterior in (4.1). The exAL augmentation preserves the Gaussian slope-path update in (3.14)–(3.15), but the scale-asymmetry blocks (σ_q, γ_q) remain non-conjugate. We therefore use two computation routes for this fixed-design posterior: an MCMC sampler with separate exAL block updates and a mean-field variational Bayes approximation with Laplace–Delta updates for the non-conjugate blocks. In the algorithms below, GIG, positive-truncated Normal, Gaussian, and inverse-gamma entries denote closed-form full conditionals for MCMC and closed-form coordinate factors for VB. The same block structure reappears in the variational updates; only the final exAL scale-asymmetry block requires a non-conjugate treatment.

4.3 Markov Chain Monte Carlo

The augmented likelihood and shrinkage hierarchy give a partially Gibbs sampler. Closed-form blocks are sampled from their named full conditionals; the complete-data targets, full conditionals,

CAVI factors, and ELBO terms for AL and exAL likelihoods under ridge and RHS priors are given in Sections S5–S9 of the supplement. Under the Nishimura–Suchard hierarchy, the global-local scales have inverse-gamma updates. Each exAL block (σ_q, γ_q) remains non-conjugate and is updated with Metropolis–Hastings or slice moves on transformed coordinates. Algorithm 1 targets (4.2); with $Q = 1$, it uses the collapsed posterior in (4.1). Ridge and AL variants omit the corresponding shrinkage, $s_{q,t}$, and γ_q blocks. Diagnostics use trace plots, effective sample size summaries, crossing diagnostics for multi-quantile grids, and monitoring of the (σ_q, γ_q) updates, including acceptance rates when Metropolis–Hastings is used.

Algorithm 1 MCMC sampler for the Q–DESN posterior. For $Q = 1$, set $\beta_{\text{base}} = \mathbf{0}_r$, place the RHS prior on $\Delta_1 = \beta_1$, and omit baseline-slope blocks. AL and ridge variants omit the corresponding $s_{q,t}$, γ_q , and shrinkage blocks.

Require: Data $\{y_t, \tilde{\mathbf{x}}_t\}_{t=1}^T$, quantile grid $p_1 < \dots < p_Q$, and priors

Require: MCMC iterations N and convergence-monitoring settings

- 1: Initialize α , β_{st} , likelihood parameters, latent variables, and RHS scales
 - 2: **if** $Q > 1$ **then**
 - 3: Initialize β_{base} and its baseline RHS scales
 - 4: **end if**
 - 5: **for** $i = 1$ to N **do**
 - 6: **Latent scales:** for each q, t , draw $v_{q,t} \mid \cdot$ from the AL or exAL GIG full conditional
 - 7: **exAL shifts:** for exAL, draw $s_{q,t} \mid \cdot$ from its positive-truncated Normal full conditional; omit this step for AL
 - 8: **Slope path:** compute $\mathbf{K}_\beta$ and $\mathbf{m}_\beta$ from (3.14)–(3.15)
 - 9: **Slope path:** draw $\beta_{\text{st}} \mid \cdot \sim \mathcal{N}(\mathbf{m}_\beta, \mathbf{K}_\beta^{-1})$
 - 10: **if** $Q > 1$ **then**
 - 11: **Baseline:** draw β_{base} from its Gaussian conditional and draw its RHS scales from inverse-gamma full conditionals
 - 12: **end if**
 - 13: **Intercepts:** draw α from its Gaussian conditional, or update the ordered-gap block when exact intercept monotonicity is imposed
 - 14: **RHS scales:** draw adjacent-difference local, global, auxiliary, and slab scales from inverse-gamma full conditionals; for $Q = 1$, this is the ordinary single-level RHS block for β_1
 - 15: **AL scale:** for AL, draw each $\sigma_q \mid \cdot$ from its inverse-gamma full conditional and omit γ_q
 - 16: **exAL scale–asymmetry:** for exAL, update each (σ_q, γ_q) on transformed coordinates targeting its non-conjugate conditional kernel
 - 17: **Monitor:** store post-burn-in draws, crossing diagnostics, and convergence summaries, with thinning only if used
 - 18: **end for**
-

4.4 Variational Approximation

For lower-cost screening, warm starts, and companion validation panels, we also use a mean-field variational approximation to the same augmented target. The Gaussian, GIG, positive-truncated Normal, inverse-gamma, and RHS scale factors retain the closed-form coordinate structure implied by (4.2). The only extra non-conjugacy introduced by the exAL likelihood is the scale-asymmetry block (σ_q, γ_q) . That block is approximated with the Laplace–Delta treatment for non-conjugate variational inference of Wang and Blei (2013), adapted to Q–DESN as detailed in the supplement. We therefore use the label VB–LD as a computational convention: it identifies the variational approximation used for screening and initialization, while the MCMC sampler remains the reference posterior computation when simulation tables report MCMC results.

Equations (4.2) and (4.1) define the posterior distributions used for inference, and Algorithm 1 implements the corresponding MCMC updates. The supplement keeps the single-level full conditionals and the full quantile-vector posterior algebra explicit: Sections S5–S8 give the single-level posterior, MCMC, VB–LD, and ELBO updates, while Section S9 gives the joint quantile-vector posterior with regularized-horseshoe shrinkage over adjacent coefficient increments, stacked Gaussian blocks, and crossing diagnostics. For a single fitted quantile, Algorithm 1 uses $Q = 1$ and the collapsed RHS prior. For multiple fitted levels, the same sampler uses the baseline and adjacent-difference blocks. Generic VB–LD factors, ELBO monitoring, the corresponding coordinate-update algorithm, and the monotone rearrangement details are given in the supplement.

5 Multi-Step Quantile Forecasting and Scoring

For a forecast origin T , a quantile path is the sequence of fitted conditional quantiles over horizons $h = 1, \dots, H$. A single-level Q–DESN fit supplies this path for one probability level p_0 . A grid of independent fits supplies one path for each fitted level $p_1 < \dots < p_K$, but the resulting grid need not be monotone at every horizon. When a monotone finite-grid quantile curve is needed for display or scoring, the default rearrangement is the weighted isotonic projection

$$\mathbf{q}_t^* = \arg \min_{q_1 \leq \dots \leq q_K} \sum_{k=1}^K w_k (\hat{q}_{k,t} - q_k)^2, \quad w_k > 0, \quad (5.1)$$

with equal weights unless an evaluation design specifies another choice in advance. The ordered grid is then linearly interpolated over $[p_1, p_K]$. Outside the fitted range, the article reports grid-level summaries unless an application explicitly states an endpoint convention. This rearrangement is a deterministic summary of fitted quantile levels. A full predictive distribution requires an additional interpolation, tail, and coupling construction.

5.1 Multi-Step Quantile Forecast Summaries

Forecast summaries are computed conditional on the fitted quantile-regression coefficients, the fixed reservoir construction, and the covariates available or specified at the forecast origin. Before forecasting, the reservoir state is obtained by running Eq. (2.1) through the observed record with observed response lags. After the forecast origin, future response lags are unknown; if they are needed for recursive multi-step forecasts, they must be generated within each Monte Carlo trajectory or supplied by a stated scenario.

For a single-level fit at probability p_0 , let $\theta_{p_0}^{(m)} = (\beta_{p_0}^{(m)}, \sigma_{p_0}^{(m)}, \gamma_{p_0}^{(m)})$ denote a posterior or variational draw. Conditional on the draw-specific future design row, the fitted quantile location is

$$\mu_{p_0, T+h}^{(m)} = \mathbf{x}_{T+h}^{(m, p_0)\top} \beta_{p_0}^{(m)}$$

for horizon h . These $\mu_{p_0, T+h}^{(m)}$ values are draws of the fitted p_0 -level quantile location. If a scalar future response path is needed for recursive lags, the fitted AL or exAL working likelihood can be used to generate

$$y_{T+h}^{(m, p_0)} \sim \text{exAL}_{p_0}(\mu_{p_0, T+h}^{(m)}, \sigma_{p_0}^{(m)}, \gamma_{p_0}^{(m)}).$$

The resulting response draws are simulations from the fitted working likelihood. Summaries of $\mu_{p_0, T+h}^{(m)}$ describe posterior or variational uncertainty in the fitted quantile location. A closed-form marginal quantile of $Y_{T+h} \mid \mathcal{F}_T$ requires an explicitly stated marginalization or simulation rule.

For the joint quantile-vector regression, a draw contains level-specific coefficients $\{\alpha_q, \beta_q, \sigma_q, \gamma_q\}_{q=1}^Q$. The same future design vector then returns the raw fitted quantile vector

$$\mu_{q, T+h}^{(m)} = \alpha_q^{(m)} + \tilde{\mathbf{x}}_{T+h}^{(m)\top} \beta_q^{(m)}, \quad q = 1, \dots, Q.$$

Let $\mathbf{q}_{T+h}^{(m)} = (\mu_{1, T+h}^{(m)}, \dots, \mu_{Q, T+h}^{(m)})^\top$. This vector estimates the conditional quantile path over the fitted grid. A scalar response predictive distribution requires the separate propagation rule below.

If exact monotonicity is required before scoring or display, let $\mathcal{M}$ denote the stated monotone rearrangement, such as isotonic projection or rearrangement (Barlow and Brunk, 1972; Chernozhukov et al., 2010), and set

$$\mathbf{q}_{T+h}^{*, (m)} = \mathcal{M}\{\mathbf{q}_{T+h}^{(m)}\}.$$

The ordered grid defines a draw-specific quantile curve $Q_{T+h}^{*, (m)}(u)$ under the stated interpolation and endpoint convention. When recursive lags require a scalar future value, the implementation uses a specified quantile-curve sampling rule, described in the supplement:

$$u_{T+h}^{(m, r)} \sim \text{Unif}(0, 1), \quad y_{T+h}^{(m, r)} = Q_{T+h}^{*, (m)}(u_{T+h}^{(m, r)}).$$

Thus the joint regression supplies the quantile curve used by the propagation rule. The scalar future response is generated from the reported quantile curve, using the stated interpolation and endpoint convention.

Single-level forecasts are scored by check loss at the fitted probability level. For a quantile grid, this article also uses aCRPS, defined as a finite-grid integrated check-loss score:

$$\text{aCRPS}(y; \mathbf{q}) = 2 \sum_{k=1}^K \omega_k \rho_{p_k}(y - q_k), \quad \rho_p(e) = e\{p - \mathbf{1}(e < 0)\}.$$

The weights ω_k are the trapezoidal quadrature weights over the fitted probability grid. Thus aCRPS approximates the CRPS identity based on integrated quantile loss over the available grid. Full continuous-distribution CRPS requires a complete quantile function and the corresponding integral. This distinction matters for the application tables below, where aCRPS compares finite sets of fitted quantile summaries.

6 Model Diagnostics and Reservoir Specification

Let

$$\mathcal{R} = \{D, \{n_d\}, \{\tilde{n}_d\}, m, T_0, \{\alpha_d\}, \{\rho_d\}, \{\pi_d^{(W)}\}, \{\pi_d^{(\text{in})}\}, f, k, \text{seed}\}$$

denote the depth, layer sizes, reduction sizes, lag memory, washout, leak rates, spectral radii, sparsity probabilities, activations, and random seed used to construct Eq. (2.1). Once $\mathcal{R}$ and the preprocessing convention are fixed, posterior inference is conditional on the resulting design matrix $\mathbf{X}$. Reservoir specification is therefore part of the empirical design: the Bayesian analysis estimates regression, likelihood, and shrinkage parameters conditional on the specified fixed feature map. Uncertainty over reservoir architectures would require an additional model layer outside the specification used here.

The ESN literature treats these choices as task-dependent design parameters. Reservoir size controls feature richness, leak rates and spectral radii jointly affect memory and stability, and input scaling determines how strongly the reservoir is driven into nonlinear regimes (Jaeger et al., 2007; Lukoševičius, 2012). For deep reservoirs, stacking layers can create hierarchical temporal representations, with additional depth justified by predictive performance and state diagnostics (Gallicchio et al., 2018). The condition $\rho_d < 1$ is a useful conservative design choice. For a driven nonlinear reservoir, empirical state diagnostics are still needed to assess stability (Yildiz et al., 2012). Several ESN modification papers identify failure modes that are visible before forecast scoring: readout coefficients can reveal task-relevant reservoir units, state feedback changes the effective recurrent dynamics through an external feedback channel, redundant state trajectories can damage readout conditioning, and random reservoir realizations can induce substantial seed-to-seed variability (Fan et al., 2017; Ehlers et al., 2025; Saadat et al., 2025; Wu et al., 2018). In Q-DESN these results enter as diagnostics, sensitivity analyses, and possible ablations. After the feature map is fixed, posterior inference is for the quantile-regression parameters conditional on that feature map.

The empirical design fixes the fitting and validation periods, held-out test period, forecast origins, horizons, quantile levels, input transformations, and scoring rules before feature maps are compared.

Numerical checks then verify input alignment, feature dimensions, washout handling, finite state values, empirical spectral radii, weakly varying or saturated reservoir states, and conditioning of the readout regression. State-matrix diagnostics include near-zero-variance columns, pairwise state correlations, singular-value or covariance spectra, effective rank, and a condition-number or eigenvalue-spread summary. These diagnostics distinguish a larger reservoir that adds nonredundant temporal features from one that mainly adds collinear columns to the fixed design matrix.

The relevant evaluation metric depends on the statistical question. In simulations with a known conditional quantile path, quantile-path RMSE, mean absolute error, bias, and correlation assess recovery of the true path. In both simulations and applications, forecast scores must use common forecast origins and horizons. For one fitted quantile level, check loss is the proper score for the reported quantile. For a quantile grid, interval score, empirical coverage, interval length, and aCRPS evaluate the finite set of fitted quantile summaries. Calibration of posterior credible bands under a misspecified working likelihood requires separate assessment.

The reservoir search used in the empirical sections varies depth, widths, lag memory, washout, leak rates, and spectral radii first; the supplement gives the full component-level checklist. Recurrent sparsity, input sparsity, input scaling, and activation choices are broadened when diagnostics indicate weak state variation, excessive saturation, or ill-conditioned design matrices. Leading specifications are compared over multiple reservoir seeds when feasible, and smaller diagnostic subsets compare VB with MCMC. The final specification is fixed before held-out evaluation. Reservoir ensembles, when used, are reported as sensitivity analyses across fixed feature maps.

Each empirical comparison is interpreted conditional on its specification grid, feature-construction defaults, state-quality summaries, forecast origins, horizons, scoring metrics, reservoir seeds, and final $\mathcal{R}$. Claims about a reservoir architecture are therefore conditional on the corresponding evaluation design. Evidence from additional seeds, forecast origins, and scoring criteria is needed before generalizing beyond the reported comparison. State-feedback or principal-neuron reinforcement methods may be useful extensions; they change the feature-generation rule and require separate validation before being combined with the Q-DESN regression.

7 Simulation Studies

The main simulation evidence is organized around two complementary questions. The first asks whether a single-quantile Q-DESN regression can recover and forecast a known dynamic quantile path under a 500-observation benchmark following the family-based comparison strategy of Yan et al. (2025). The target quantile level is fixed in advance, the data-generating mechanism is known, and competing models are compared under matched error families, fitting-sample sizes, and quantile levels. The second asks whether a joint multi-quantile regression can recover several known conditional quantile paths under posterior simulation and maintain coherent raw quantile grids over a fixed set of levels. Companion VB screening and forecast-comparison tables are kept in the supplement rather than used as the main article evidence.

In contrast to the linear-regression setting of Yan et al. (2025), the oracle quantile paths are dynamic, and Q-DESN represents those paths through high-dimensional deterministic reservoir features. The design therefore separates two questions: how accurately a fitted model recovers the known conditional quantile path, and how well it forecasts held-out dynamic blocks. We write p for the single target level in the first study and τ for levels in the multi-quantile grid.

The main-text tables report MCMC summaries under fixed evaluation designs: data-generating mechanisms, rolling-origin definitions, fitting samples, forecast-validation periods, scoring rules, and model classes are specified before the reported comparisons are formed. The supplement gives VB companion panels and additional diagnostic summaries for the same simulation settings. The single-quantile and multi-quantile studies use separate fixed designs, so their results are interpreted as complementary evidence rather than as one pooled experiment.

7.1 Single-Quantile Dynamic Data-Generating Processes

For each error family f and target quantile level $p \in \{0.05, 0.25, 0.50\}$, the synthetic series is generated as

$$y_{t,p,f} = \mu_{t,f} + \varepsilon_{t,p,f}, \quad \varepsilon_{t,p,f} = e_{t,f} - H_f^{-1}(p),$$

where H_f is the distribution function of the raw error $e_{t,f}$. The shift by $H_f^{-1}(p)$ makes the p -quantile of $\varepsilon_{t,p,f}$ equal to zero. Hence

$$Q_p(y_{t,p,f} \mid \boldsymbol{\theta}_{t,f}) = \mu_{t,f}.$$

The oracle quantile path is generated by a six-dimensional dynamic linear state with local level, local slope, and two harmonic pairs:

$$\mu_{t,f} = \mathbf{F}^\top \boldsymbol{\theta}_{t,f}, \quad \mathbf{F} = (1, 0, 1, 0, 1, 0)^\top,$$

with local linear trend and two sine-cosine harmonic pairs. The seasonal period is 90, the harmonics are 1 and 2, and the state innovation covariance $\mathbf{W}_\theta$ has diagonal square roots

$$\sqrt{\text{diag}(\mathbf{W}_\theta)} = (0.005, 0.00002, 0.004, 0.004, 0.003, 0.003).$$

The initial covariance is $\mathbf{C}_0 = 0.01\mathbf{I}_6$, and each simulated trajectory is initialized deterministically at the family-specific state in Table 1. For a fixed family, the same latent path $\{\mu_{t,f}\}$ is used across the three values of p ; only the centering constant $H_f^{-1}(p)$ changes.

7.2 Single-Quantile Fit and Forecast Protocol

Each synthetic series is first run for $T_{\text{warm}} = 2000$ warmup observations, which are discarded before analysis. The next $T_{\text{main}} = 10000$ observations form the retained series used for fitting and validation. Models are fit at the end of the first 9000 retained observations, and the final 1000 retained observations are held out for forecast evaluation. We report two fitting-sample sizes,

Family	Raw error distribution H_f	Initial level	Initial slope	Seasonal settings
Gaussian	$\mathcal{N}(0, 10^2)$	40	0.012	harmonics (24, 0.35), (8, -0.80)
Laplace	Laplace with scale 10	35	0.011	harmonics (28, -0.15), (10, 0.75)
Gaussian mixture	$0.1\mathcal{N}(0, 0.5^2) + 0.9\mathcal{N}(1, 15^2)$	45	0.014	harmonics (32, 0.85), (12, -1.25)

Table 1: Single-quantile dynamic data-generating mechanisms. Raw errors are shifted by $H_f^{-1}(p)$ so that the p -quantile of the observation error is zero; the same latent dynamic quantile path is then shared across target levels within each family. Seasonal entries report amplitude and phase for the two harmonic pairs in the initial state.

$T_{\text{fit}} = 500$ and $T_{\text{fit}} = 5000$, using respectively the 500 or 5000 observations immediately preceding the fitting origin. Forecast scores are computed from rolling origins within the held-out period; origins are spaced 30 observations apart, and each origin is scored for horizons up to 30 steps ahead. The same protocol is used for all model specifications and score calculations.

7.3 Single-Quantile Competing Methods

For each family and target quantile level in the single-quantile simulation comparison, the reported Q-DESN rows vary the working likelihood and inference method while using regularized-horseshoe (RHS) coefficient shrinkage. Entries labeled Q-DESN AL-RHS use the AL working likelihood, whereas entries labeled Q-DESN exAL-RHS use the exAL working likelihood. Historical ridge fits are omitted from the main comparison because the reported comparison focuses on Bayesian shrinkage fits under the fitting-sample preprocessing rule. The reservoir specification and random seed are fixed for every displayed metric. Dynamic quantile linear model (DQLM) and extended dynamic quantile linear model (exDQLM) baselines are fit to the same family, quantile level, fitting-sample size, and forecast-evaluation protocol using `exdqim` (Barata et al., 2026).

The primary comparison is a fixed-design benchmark against structured dynamic quantile linear models under linear dynamic mechanisms that favor those baselines. Repeated data-generating realizations, broader reservoir tuning, and non-Bayesian ESN regression baselines are left to the supplementary diagnostics and future sensitivity analyses.

7.4 Single-Quantile Criteria for Comparison

Let $\hat{q}_t$ denote a posterior or variational summary of the fitted conditional p -quantile, such as the posterior median of the linear-predictor location. Because the oracle path μ_t is known, fitting-sample recovery can be evaluated using quantile mean absolute error, quantile RMSE, bias, correlation, and check loss at the target level p . Forecast accuracy is evaluated on held-out outcomes under common forecast origins and horizons for every model. The reported tables show fit recovery and forecast scoring.

For single-level quantile fits, the primary proper score is check loss at the fitted level. Quantile MAE and quantile RMSE compare the fitted or forecast quantile summary directly with the known oracle quantile path, while check loss scores realized observations relative to the reported quantile. These metrics therefore answer different questions and are reported separately.

Because the simulation uses one fixed data-generating realization for each family and quantile level, the tables report deterministic criterion values for that data-generating realization rather than Monte Carlo means over repeated data sets. Any boldface or ranking in the displayed MCMC tables is interpreted within a fixed family, quantile level, fitting-sample size, model family, and metric. The displayed Q-DESN values form a metric-specific summary: each entry is the best finite value observed for that model, family, quantile level, and metric. Consequently, entries within one model row may originate from different model specifications. The supplementary tables identify the corresponding diagnostic status.

7.5 Single-Quantile Fit and Forecast Results

The main comparison uses the shorter fitting design: 500 observations before the fitting origin, followed by 1000 held-out observations for forecast scoring. Forecast scores use rolling origins spaced 30 observations apart and horizons up to 30 steps ahead, at the three target levels $p = 0.05, 0.25, 0.50$. Tables 2–4 report the MCMC family-specific metric summaries. The entries are computed from horizon-level rolling-origin comparisons: fit RMSE measures recovery of the known oracle quantile path on the fitting sample, forecast MAE measures oracle-quantile forecast error averaged across scored origin-horizon pairs, and forecast check loss scores realized observations at the fitted quantile level. Boldface marks the lowest displayed MCMC value within each quantile level and metric. Diagnostic status is disclosed alongside the scores: daggers mark rows with a warning in the supporting diagnostic summary, and double daggers mark rows with a failed diagnostic check.

All displayed Q-DESN and exQ-DESN values use feature scaling and other learned preprocessing quantities estimated from observations available through the fitting origin. The DQLM and exDQLM baselines use the same source trajectories, fitting sample, forecast origins, horizons, and scoring rules. The exAL-RHS MCMC rows use the v -augmented transition described in the supplement, which separates the scale update from the shape update and samples the shape on the logit transform of its bounded support.

Two case-specific repeated-chain confirmations yielded four further strict forecast improvements. For Gaussian $p = 0.05$, AL-RHS reduced forecast MAE from 8.410 to 6.917 and forecast check loss from 1.221 to 1.200. For the Gaussian-mixture family at $p = 0.50$, exAL-RHS reduced forecast MAE from 2.562 to 1.420 and forecast check loss from 5.610 to 5.487. These replacements are three-chain means. Fit RMSE and all other displayed criteria retain their previous metric-specific sources, and diagnostic grades remain disclosed in the supporting record.

The supplement gives VB companion panels for the same source trajectories and a five-chain sensitivity analysis for the Gaussian family at $p = 0.25$. The sensitivity analysis assesses whether the reported Q-DESN summaries persist across independently seeded chains using a common posterior

Model	$p = 0.05$			$p = 0.25$			$p = 0.50$		
	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss
DQLM	3.036	3.829	1.103	3.811	2.209	3.328	2.590	1.161	4.029
exDQLM	1.935	16.23	4.313	2.344	6.634	4.132	2.808	1.637	4.109
Q-DESN AL-RHS [†]	2.779	6.917	1.200	2.176	2.191	3.292	1.523	2.104	4.073
Q-DESN exAL-RHS [‡]	2.149	2.863	1.080	1.366	2.415	3.317	1.128	1.971	4.059

Table 2: MCMC single-quantile fit-and-forecast comparison for the Gaussian simulation family. Q-DESN entries are reported within each model, target level, and criterion under the fixed evaluation design. When repeated-chain summaries are available, the displayed value uses the corresponding repeated-chain average. Q-DESN preprocessing is estimated only from the fitting sample. Forecast criteria average over the scored origin-horizon pairs in the 1000-observation held-out period, with origins spaced 30 observations apart and horizons up to 30 steps ahead. Lower values are better, and boldface marks the best displayed value within each target level and criterion. A dagger marks a warning in the supporting diagnostics; a double dagger marks a failed diagnostic check.

Model	$p = 0.05$			$p = 0.25$			$p = 0.50$		
	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss
DQLM	3.663	9.791	1.968	2.850	3.520	4.548	1.774	1.337	5.082
exDQLM	5.677	22.44	5.154	1.710	6.811	5.106	1.774	2.013	5.206
Q-DESN AL-RHS [†]	5.420	3.893	1.873	2.256	1.434	4.379	1.353	0.7553	5.008
Q-DESN exAL-RHS [‡]	6.409	2.362	1.855	1.383	0.6696	4.343	1.075	0.6772	5.008

Table 3: MCMC single-quantile fit-and-forecast comparison for the Laplace simulation family. Q-DESN entries are reported within each model, target level, and criterion under the fixed evaluation design. When repeated-chain summaries are available, the displayed value uses the corresponding repeated-chain average. Q-DESN preprocessing is estimated only from the fitting sample. Forecast criteria average over the scored origin-horizon pairs in the 1000-observation held-out period, with origins spaced 30 observations apart and horizons up to 30 steps ahead. Lower values are better, and boldface marks the best displayed value within each target level and criterion. A dagger marks a warning in the supporting diagnostics; a double dagger marks a failed diagnostic check.

Model	$p = 0.05$			$p = 0.25$			$p = 0.50$		
	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss	Fit RMSE	Forecast MAE	Forecast check loss
DQLM	2.623	3.003	1.505	4.740	2.487	4.540	2.226	1.932	5.556
exDQLM	2.554	23.91	5.857	1.381	8.941	5.649	2.408	2.207	5.607
Q-DESN AL-RHS [†]	3.154	3.006	1.564	1.855	0.9122	4.467	1.271	0.9074	5.441
Q-DESN exAL-RHS [‡]	3.354	2.913	1.527	1.380	1.819	4.488	0.9177	1.420	5.487

Table 4: MCMC single-quantile fit-and-forecast comparison for the Gaussian mixture simulation family. Q-DESN entries are reported within each model, target level, and criterion under the fixed evaluation design. When repeated-chain summaries are available, the displayed value uses the corresponding repeated-chain average. Q-DESN preprocessing is estimated only from the fitting sample. Forecast criteria average over the scored origin-horizon pairs in the 1000-observation held-out period, with origins spaced 30 observations apart and horizons up to 30 steps ahead. Lower values are better, and boldface marks the best displayed value within each target level and criterion. A dagger marks a warning in the supporting diagnostics; a double dagger marks a failed diagnostic check.

point-path summary.

7.6 Joint Multi-Quantile Study with a Pre-Specified Feature Design

The second simulation study assesses whether a shared quantile-vector regression can recover several conditional quantile paths simultaneously under a feature design fixed before evaluation. Reservoir and prior specifications are chosen separately for each mechanism and model using VB or VB-LD. A replicated robustness experiment fits the 32 scenario–model specifications on 50 independently generated data-generating realizations per comparison, for 1,600 variational fits in total. The main comparison is a balanced MCMC analysis of the same four regression–likelihood combinations: Joint QDESN RHS, Independent QDESN RHS, Joint exQDESN RHS, and Independent exQDESN RHS. Each MCMC run is initialized from its corresponding variational fit and uses the scenario-specific specification fixed before sampling. The comparison is restricted to posterior quantile-grid paths under the stated working likelihoods.

The exAL rows include eleven scale-collapsed fits whose quantile summaries satisfy the stability criterion and five stable exAL specifications from the same scenario-level evaluation design. Mixing qualifications are reported in the diagnostic summaries. Across the scale-collapsed exAL fits, fitting-sample MAE improves consistently, while forecast MAE improves in five of the sixteen exAL scenario–model comparisons.

Each regression is fit across

$$\tau \in \{0.05, 0.10, 0.25, 0.50, 0.75, 0.90, 0.95\}.$$

The balanced MCMC comparison grid uses eight synthetic data-generating mechanisms, a pre-specified feature design, and a 500-observation fitting sample. The three bridge mechanisms use Gaussian, Laplace, and Gaussian-mixture innovations. The five stress mechanisms introduce asymmetric tails, persistent heavy tails, Student- t errors, regime shifts, and nonlinear reservoir-friendly dynamics. The resulting grid contains 32 scenario–model MCMC rows. Fitted and forecast quantiles are compared with the known conditional quantile paths; realized observations are scored by check loss and finite-grid aCRPS; and hit-rate errors summarize marginal calibration. Reported scores use the pre-specified monotone rearrangement, while raw adjacent-level crossings before that rearrangement are retained as numerical coherence diagnostics.

Scenario-level results. Table 5 summarizes the balanced MCMC analysis at the scenario level. Entries report forecast MAE against the oracle quantile grid, with fitting-sample MAE and raw held-out crossing counts in parentheses; boldface marks the lowest within-scenario forecast MAE. Joint QDESN RHS is lowest in the Laplace, nonlinear, and normal mechanisms. Independent QDESN RHS is lowest in the asymmetric-tail, regime-shift, and Student- t location-scale mechanisms, while Independent exQDESN RHS is lowest in the Gaussian-mixture and persistent-heavy-tail mechanisms. Joint exQDESN RHS has larger forecast MAE than at least one competing structure in each mechanism. These numerical methods with the lowest displayed scores are descriptive because

Scenario	Mechanism	Joint QDES AL-RHS	Independent QDES AL-RHS	Joint exQDES exAL-RHS	Independent exQDES exAL-RHS
Asymmetric Laplace Tail	Asymmetric Laplace; seasonal AR(1) location-scale	0.099 (0.084; 1)	0.083 (0.088; 7)	0.093 (0.081; 0)	0.089 (0.077; 0)
Gaussian-Mixture Bridge	Gaussian mixture; seasonal AR(1) location-scale	0.096 (0.097; 0)	0.099 (0.101; 0)	0.106 (0.096; 0)	0.094 (0.095; 0)
Laplace Bridge	Laplace; seasonal AR(1) location-scale	0.085 (0.089; 0)	0.094 (0.091; 7)	0.111 (0.113; 0)	0.106 (0.117; 0)
Nonlinear Reservoir Friendly	Gaussian mixture; nonlinear reservoir dynamics	0.104 (0.109; 0)	0.116 (0.112; 1)	0.157 (0.151; 0)	0.149 (0.143; 0)
Normal Bridge	Gaussian; seasonal AR(1) location-scale	0.081 (0.082; 0)	0.082 (0.082; 2)	0.115 (0.097; 0)	0.104 (0.101; 0)
Persistent Heavy Tail	Student-t; persistent seasonal AR(1) dynamics	0.131 (0.120; 0)	0.124 (0.118; 0)	0.130 (0.103; 0)	0.106 (0.100; 0)
Regime Shift	Student-t; regime-shift location-scale	0.153 (0.085; 0)	0.096 (0.091; 6)	0.185 (0.090; 0)	0.113 (0.087; 0)
Student-t Location-Scale	Student-t; seasonal AR(1) location-scale	0.080 (0.062; 0)	0.071 (0.065; 1)	0.116 (0.093; 0)	0.099 (0.092; 0)

Table 5: Scenario-level MCMC evaluation for the joint multi-quantile simulation study. Entries are forecast MAE with fitting-sample MAE and the number of raw held-out adjacent-level crossings in parentheses, written as forecast MAE (fitting-sample MAE; raw crossings). Boldface marks the lowest forecast MAE within each pre-specified eight-mechanism scenario. QDES uses the AL working likelihood, exQDES uses the exAL working likelihood, and RHS denotes the regularized horseshoe prior. All scores use the monotone quantile grid; raw crossings are retained only as pre-rearrangement diagnostics. exAL rows are included when finite-path and quantile-stability criteria are satisfied under the balanced evaluation design.

complete best-versus-runner-up Monte Carlo error is unavailable for all comparisons.

Coherence diagnostics and interpretation. Fit recovery gives a complementary result: Joint QDES RHS has the lowest MAE in five of the eight fitting samples, with Independent exQDES RHS lowest in the three remaining cases. Quantile coherence also distinguishes the regression structures. Independent QDES RHS accounts for 24 of the 25 raw held-out crossings; Joint QDES RHS accounts for one, and both exAL rows account for none. All 32 rows have zero crossings after the monotone rearrangement. The replicated VB experiment further gives lower median paired forecast MAE to AL than exAL for both regression structures in every mechanism, although the frequency and size of the differences vary by scenario. Together, these results identify Joint QDES RHS as the principal joint-regression reference in this design. The exAL shape parameter is useful in some comparisons, but its forecast quantile-path accuracy is heterogeneous across mechanisms. The supplement reports the MCMC settings, crossing diagnostics, metric-specific lowest displayed scores, and replicated AL–exAL contrasts.

8 Application: GloFAS Retrospective Streamflow Quantile Forecasting

Retrospective case study. The empirical application considers a medium-range retrospective streamflow quantile-forecasting case using the Global Flood Awareness System (GloFAS) and a reference gauge record. GloFAS is a global hydrological forecast and monitoring system jointly developed by the European Commission and the European Centre for Medium-Range Weather Forecasts (ECMWF) and operated within the Copernicus Emergency Management Service (Alfieri

et al., 2013; Copernicus Emergency Management Service, 2026). Its operational forecast and reforecast products provide ensemble river-discharge information for flood early warning and forecast verification (Harrigan et al., 2023). Previous calibration work has focused on hydrological model parameters using daily streamflow observations (Hirpa et al., 2018); here Q-DESN is used as a statistical quantile-regression correction for one forecast system and one specified forecast origin. The case is retrospective: the reported specification uses a blended future-covariate information set and is evaluated at a single origin.

8.1 Data and Forecast-Origin Design

Let T denote a forecast origin and let H be the issued forecast horizon. The historical reference record is y_t , the transformed USGS streamflow through time T . The retrospective GloFAS product over the same period is g_t^{ret} , and $g_{T+h,m}^{\text{ens}}$ denotes GloFAS ensemble member m issued at origin T for horizon $h = 1, \dots, H$. The response values $y_{T+1}, \dots, y_{T+H}$ are held out during prediction and used later for validation.

Retrospective information set. The reported analysis uses the forecast origin 2022-12-25 and scores seven fitted quantile levels over 28 issued target dates. The application is a single-origin retrospective case. Its response values are held out for scoring, and the covariate policy includes realized-history and blended Global Ensemble Forecast System (GEFS) forecast information with realized-future weather contributions for precipitation and soil-moisture covariates. This restriction makes the example a controlled retrospective comparison before broader operational evaluations are attempted.

8.2 Application Model

Persistence-anchored discrepancy. For $t \leq T$, the application uses two deterministic Q-DESN feature maps. The reference block $\mathbf{x}_t^{\text{ref}}$ is driven by the transformed reference history and the same covariate policy. The discrepancy block $\mathbf{x}_t^{\text{disc}}$ is driven by the retrospective discrepancy $g_t^{\text{ret}} - y_t$, using the same preprocessing, lag, reservoir, reducer, and washout conventions. For $t > T$, the reported retrospective specification uses the persistence-anchored innovation rule: the future discrepancy feature is anchored at the last historical retrospective discrepancy and the discrepancy regression models an adjustment around that persistence baseline. The empirical issued-ensemble quantile is retained as a raw GloFAS comparator and as part of the forecast-system evidence table; the future ESN feature map for the discrepancy is driven by the persistence-anchored discrepancy state.

For a fixed quantile level p_0 , the reference-process quantile is represented by one Q-DESN regression and the GloFAS discrepancy by a second regression,

$$q_{p_0,t}^{\text{ref}} = \mathbf{x}_t^{\text{ref}\top} \boldsymbol{\beta}_{p_0}, \quad d_{p_0,t}^{\text{glo}} = \mathbf{x}_t^{\text{disc}\top} \boldsymbol{\alpha}_{p_0}, \quad q_{p_0,t}^{\text{glo}} = q_{p_0,t}^{\text{ref}} + d_{p_0,t}^{\text{glo}}.$$

The path $q_{p_0,t}^{\text{ref}}$ is the reference-process quantile of interest, while $d_{p_0,t}^{\text{glo}}$ is the quantile-specific

discrepancy between GloFAS and the reference process. The retrospective and ensemble GloFAS products inform the forecast-system quantile path $q_{p_0,t}^{\text{glo}}$; the scored reference quantity remains $q_{p_0,T+h}^{\text{ref}}$.

At forecast origin T , the issued GloFAS ensemble supplies river-discharge alternatives from the forecast system at each covered target date. The reported application summary is a VB quantile-regression summary of the reference process and the persistence-anchored discrepancy adjustment. For variational draw or summary index $s = 1, \dots, S$, the construction records the reference quantile and the resulting forecast-system quantile:

$$q_{p_0,T+h}^{\text{ref},(s)} = \mathbf{x}_{T+h}^{\text{ref},(s)\top} \boldsymbol{\beta}_{p_0}^{(s)}, \quad q_{p_0,T+h}^{\text{glo},(s)} = q_{p_0,T+h}^{\text{ref},(s)} + d_T^{\text{ret}} + \mathbf{x}_{T+h}^{\text{disc},(s)\top} \boldsymbol{\alpha}_{p_0}^{(s)}.$$

Here d_T^{ret} is the last historical retrospective discrepancy on the transformed scale. Posterior means, medians, intervals, and monotone multi-quantile summaries are computed after this quantile construction.

The working-likelihood specification is

$$\begin{aligned} y_t \mid q_{p_0,t}^{\text{ref}}, \sigma_{\text{ref}}, \gamma_{\text{ref}} &\sim \text{exAL}_{p_0}(q_{p_0,t}^{\text{ref}}, \sigma_{\text{ref}}, \gamma_{\text{ref}}), \\ g_t^{\text{ret}} \mid q_{p_0,t}^{\text{glo}}, \sigma_{\text{glo}}, \gamma_{\text{glo}} &\sim \text{exAL}_{p_0}(q_{p_0,t}^{\text{glo}}, \sigma_{\text{glo}}, \gamma_{\text{glo}}), \\ g_{T+h,m}^{\text{ens}} \mid q_{p_0,T+h}^{\text{glo}}, \sigma_{\text{glo}}, \gamma_{\text{glo}} &\sim \text{exAL}_{p_0}(q_{p_0,T+h}^{\text{glo}}, \sigma_{\text{glo}}, \gamma_{\text{glo}}), \quad h = 1, \dots, H, \quad m = 1, \dots, M. \end{aligned}$$

This display gives the general exAL form. The AL specialization used by the reported AL analysis omits s -latent variables and γ blocks, and replaces the exAL constants by the corresponding AL constants at the fitted level. The issued ensemble members enter through a working likelihood that treats them as conditionally independent and exchangeable given the forecast-system quantile path and source-specific likelihood parameters. ECMWF ensemble forecasts use perturbed initial states and model perturbations to produce distinct forecast alternatives (European Centre for Medium-Range Weather Forecasts, 2026); the conditional-independence factorization is therefore a modeling approximation for the Q-DESN regression layer. In exAL variants, the default application model shares σ_{glo} and γ_{glo} between retrospective and issued-ensemble GloFAS rows; under the AL specialization, only the corresponding σ_{glo} block is shared.

For this retrospective analysis, the seven fitted quantile levels are 0.05, 0.15, 0.35, 0.50, 0.65, 0.80, and 0.95. They are fitted independently with the AL special case and the VB approximation. The streamflow application is intended as a fast retrospective quantile-regression analysis, so the reported run uses Q-DESN under AL with independent fitting and monotone rearrangement. The variational quantile summaries are passed through the monotone rearrangement to form the displayed intervals. The displayed bands use the fitted quantile models directly, followed only by monotone rearrangement; the reservoir features, likelihood, priors, and fitted regression coefficients remain those of the reported fit.

The reported fit uses separate DESN feature maps for the reference quantile path and the GloFAS discrepancy path. Both maps have depth $D = 1$, layer widths (300), reducer widths (*none*), memory

360, washout 500, spectral radii (0.95), recurrent sparsity (0.03), and input inclusion probabilities (1). The reference map uses leak rates (0.1) and global/bias input scales 0.18/0.18, and reservoir seed 20260512; the discrepancy map uses leak rates (0.1), global/bias input scales 0.18/0.18, and reservoir seed 20261521. Forecast-window discrepancy states use the persistence-anchored innovation transition.

The two coefficient blocks use separate regularized-horseshoe global scales, $\tau_{0,\text{ref}} = 0.1$ and $\tau_{0,\text{disc}} = 0.001$, respectively.

8.3 Retrospective Application Results

The reported fit was evaluated over 12,495 observed dates through the forecast origin. After monotone rearrangement, its observational-window aCRPS was 0.0445, and empirical 90% interval coverage was 0.906. This evidence governs the specification choice; the forecast period remains a held-out assessment. The retrospective diagnostic evaluation also identified occasional extreme fitted values in the $p_0 = 0.95$ path, so the upper-tail component remains subject to targeted sensitivity analysis.

Forecast-period scores. Table 6 reports the forecast scores for the reported retrospective analysis. Relative to raw GloFAS, Q-DESN reduces the mean check loss from 0.7639 to 0.5654, the interval score from 25.2028 to 10.8032, and aCRPS from 1.4424 to 1.1319. Mean empirical interval coverage increases from 0.000 to 0.357, but the Q-DESN value remains far below the nominal 90% level. These summaries describe one specified forecast origin under a blended retrospective information set. Multi-origin operational validation and calibrated predictive intervals require broader experiments.

Table 6: Forecast scoring for the GloFAS retrospective analysis. Scores are computed on the transformed streamflow scale for the forecast origin 2022-12-25. Lower check loss, interval score, and aCRPS are better; aCRPS is the finite-grid trapezoidal integrated check-loss score over the seven fitted levels. Mean coverage reports empirical coverage for the nominal 90% reported interval and is interpreted descriptively. The final column gives the relative reduction in mean check loss.

Model	Horizons	Check	Interval	aCRPS	Coverage	Check reduction
Q-DESN correction	28	0.5654	10.8032	1.1319	0.357	26.0%
Raw GloFAS	28	0.7639	25.2028	1.4424	0.000	Reference




Figure 1: Synthesized Q-DESN quantile bands for the GloFAS retrospective analysis over the final 60 observed dates and the following 28-day forecast period. The nested blue bands are the 90%, 65%, and 30% quantile intervals, and the blue path is the median. The vertical dashed line marks the forecast origin. Black and red paths denote observed and held-out USGS streamflow, respectively; orange denotes retrospective GloFAS, and thin gray paths show all 51 issued GloFAS ensemble members, with their median dashed. Forecast scores use the held-out measurements in Table 6.

9 Application: PriceFM Retrospective Electricity-Price Comparison

As a second empirical case, we consider a retrospective comparison of probabilistic day-ahead electricity-price forecasts using version 4 of the PriceFM benchmark (Yu et al., 2026). The current benchmark release spans 1 January 2022 through 1 January 2026. The scored quantity is original-unit quantile loss on the PriceFM paper quantile grid under fold-aligned test timestamps. The released data cover European bidding regions at 15-minute resolution and include day-ahead price, load, solar, and wind forecasts. PriceFM itself is a graph-informed foundation model: it uses region-level exogenous information and a sparse graph mask derived from the European transmission topology. The reported comparison is a 114 region-fold comparison against fold-aligned PriceFM predictions. It evaluates the aligned benchmark panel used in this article.

9.1 Comparison Design

The reported benchmark contains 114 region/fold evaluations, spanning 38 European bidding regions and 3 rolling folds. Each evaluation is scored on the same fold-aligned test timestamps used by the PriceFM predictions. The scoring metric is original-unit average quantile loss (AQL) on the PriceFM paper quantile grid 0.10, 0.25, 0.45, 0.50, 0.55, 0.75, 0.90. Throughout this section, Δ AQL denotes Q-DESN AQL minus PriceFM AQL. Negative values indicate lower Q-DESN AQL; positive values indicate lower PriceFM AQL.

Retrospective covariates. The benchmark panel is assembled from reported Q-DESN forecasts and PriceFM paper-quantile forecasts. The comparison uses the existing PriceFM predictions. Across the benchmark, 58 reported region/fold evaluations use own-region predictors and 56 reported region/fold evaluations add PriceFM neighborhood summaries. Both input policies use retrospectively observed own-region load, solar, and wind lead covariates in the comparison; evaluations with neighborhood summaries add retrospectively observed neighboring-region lead summaries. The score comparison is therefore retrospective and fold-aligned. The benchmark is interpreted as an application of the reported Q-DESN quantile-regression analyses. A PriceFM study of the joint quantile-vector RHS prior would require the same fold-aligned design with that joint regression fitted directly.

PriceFM reports two model-comparison protocols. Its Table II evaluates an all-region fitted setting, whereas its Table III evaluates leave-one-region-out generalization. Because the Q-DESN specifications here are fitted with observations from each target region, Table II is the relevant paper-level analogue. Ranking for the present region-specific fits is therefore based on that comparison.

9.2 Full-Surface Results

Across the aligned benchmark, Q-DESN has lower AQL in 66 of 114 region/fold evaluations (57.9%), 12 evaluations are near-equivalent under the pre-specified near-equivalence threshold, and PriceFM has lower AQL in 36 evaluations. The mean Q-DESN AQL is 6.826, compared with mean PriceFM AQL 7.039, giving mean Δ AQL -0.213 and median Δ AQL -0.083. The aggregate difference is therefore modest and heterogeneous. The fold-level supplementary summary shows that the mean Δ AQL is most negative in fold 2, modestly negative in fold 1, and essentially neutral in fold 3. On the same 114-comparison panel, the region-fold macro mean AQCR is 2.64% for Q-DESN and 0.00% for the released PriceFM checkpoint. Mean MAE is 16.728 versus 17.274, and mean RMSE is 25.453 versus 26.336, respectively. Thus Q-DESN reduces aligned AQL, MAE, and RMSE by approximately 3.0%, 3.2%, and 3.4%, while retaining a nonzero quantile-crossing rate.

The directly aligned checkpoint is the appropriate distribution-level comparator because its predictions are aligned row by row with the Q-DESN predictions. Its AQL differs from the PriceFM paper’s optimized Table II AQL, so the two PriceFM rows are kept separate. Numerically, the aligned Q-DESN AQL lies within the range reported for the graph baselines in PriceFM Table II, but differences in fitted objects and aggregation preclude a cross-panel rank claim. The fold-level win, tie, and loss summary is reported as a supplementary diagnostic.

Table 7: PriceFM-aligned retrospective model comparison. The first panel reports direct region–fold macro means from the common 38-region, three-fold comparison; boldface marks the better of those two directly aligned rows. The second panel reproduces the Table II values from PriceFM version 4 (Yu et al., 2026). The paper-reported panel uses a distinct fitted-model and summary construction and is therefore contextual for the aligned comparison panel. AQL, MAE, and RMSE are in EUR/MWh; AQCR is in percent. Rank is reproduced only for the paper-reported panel; dashes denote quantities without a cross-panel rank. Lower is better. The comparison uses retrospectively observed own-region covariate leads in all reported rows and additional retrospectively observed neighboring-region leads when neighborhood summaries are included.

Model	Probabilistic		Pointwise		Rank
	AQL	AQCR (%)	MAE	RMSE	
<i>Direct fold-aligned comparison (this article; 114 region/fold evaluations)</i>					
Reported Q–DESN	6.83	2.64	16.73	25.45	—
Fixed PriceFM predictions	7.04	0.00	17.27	26.34	—
<i>PriceFM v4 Table II (paper-reported all-region fitted evaluation)</i>					
Naïve ¹	15.29	0.00	22.06	34.68	11
Naïve ²	15.35	0.00	23.31	34.31	13
Naïve ³	15.46	0.00	22.64	32.61	12
FEDFormer	8.22	15.33	20.15	31.75	8
PatchTST	8.06	18.21	20.20	31.59	8
iTransformer	8.24	13.96	21.03	32.11	9
TimesNet	7.98	13.42	19.48	30.94	7
TimeXer	8.30	14.77	21.94	31.88	10
GraphConv	6.61	6.88	16.81	25.97	4
GraphAttn	7.13	10.33	17.00	26.11	6
GraphSAGE	6.78	6.01	17.56	26.03	5
GraphDiffusion	6.69	5.72	16.44	25.93	2
GraphARMA	6.72	6.03	16.56	25.84	3
PriceFM	5.80	0.00	14.28	22.39	1

The input set also matters. Evaluations with neighborhood summaries have the most negative average Δ AQL, whereas own-region-only evaluations are closer to neutral. Because input sets and fitted model specifications vary across region/fold evaluations, the input-set comparison is descriptive. It summarizes where the reported Q-DESN analyses are strongest on the aligned benchmark. Input-set and horizon-block breakdowns are reported in the supplement.

9.3 Horizon Diagnostics and Interpretation

Horizon-block diagnostics are available for 72 of the 114 region/fold evaluations. Negative Δ values indicate lower Q-DESN AQL. Because the horizon summaries cover a retained subset of the aligned panel, they are interpreted descriptively.

The resulting evidence supports a conservative interpretation. In this retrospective comparison, Q-DESN is a competitive fixed-reservoir quantile-regression baseline against fold-aligned PriceFM predictions, with lower mean and median AQL and lower region-fold macro MAE and RMSE in aggregate. This advantage is descriptive, relative to the fixed PriceFM prediction file used in the aligned panel, and heterogeneous across folds and regions. The separately reported optimized PriceFM Table II result is contextual, and paired uncertainty intervals are left for a future benchmark analysis. A new PriceFM comparison would require a stated change to the information set, model class, comparison rule, or joint quantile-regression design. A fully joint PriceFM run would be needed before making application-level claims about the joint multi-quantile regression model.

10 Discussion

Q-DESN treats the fixed DESN reservoir as a nonlinear feature map and places Bayesian quantile inference on the regression coefficients. Posterior simulation and VB update the coefficients, likelihood, and shrinkage parameters, while the recurrent weights, reducers, scaling constants, and washout convention remain fixed. The empirical results evaluate this conditional Bayesian quantile-regression analysis under specified reservoir constructions.

The likelihood is also deliberately limited in interpretation. The AL and exAL families are used as quantile-targeting working likelihoods. They provide posterior distributions and conditionally tractable augmented updates. The simulation and application sections report fit recovery, forecast scoring, empirical predictive coverage, monotonicity, and convergence diagnostics to assess the model-based summaries under possible misspecification. Sandwich adjustments and other general-Bayes calibration checks remain important sensitivity directions, especially for composite multi-quantile likelihoods.

The multi-quantile construction makes the same distinction. Independent single-level Q-DESN fits followed by isotonic regression or rearrangement produce monotone quantile curves from level-wise regressions. The joint quantile-vector extension shares information across quantile levels through RHS shrinkage on adjacent coefficient differences, giving a soft non-crossing mechanism. Exact monotonicity for downstream use requires an ordered-intercept parameterization, projection,

rearrangement, or another explicitly specified monotone reporting rule.

The empirical evidence supports conditional claims. The single-quantile dynamic simulations test recovery and held-out forecasting under known data-generating mechanisms. The joint multi-quantile synthetic study checks coherence under known conditional quantile paths. The GloFAS application studies one retrospective streamflow forecast origin with independent Q-DESN AL-VB fits, blended future weather covariates, monotone reporting, and undercoverage of the nominal 90% interval. The PriceFM comparison is fold-aligned against stored benchmark predictions with retrospectively observed lead covariates. Broader rankings of Q-DESN against dynamic quantile or foundation-model alternatives require additional multi-design comparisons.

Several extensions are natural. On the statistical side, future work should study calibrated composite-likelihood scaling, hard or probabilistic non-crossing constraints, richer covariance structures for quantile differences, and sensitivity to reservoir scaling on bounded feature domains. On the computational side, the joint model calls for sparse precision solvers, targeted covariance extraction for VB, and diagnostics comparing VB with MCMC on small joint problems. On the application side, deployable multi-origin GloFAS studies with origin-available covariates, paired PriceFM uncertainty analyses, and explicitly joint PriceFM runs are needed before the joint quantile-vector prior can be interpreted in an application-level benchmark.

Within these boundaries, the article establishes Q-DESN as a fixed-feature Bayesian quantile-regression framework whose inference, monotone rearrangement, diagnostics, and empirical claims are tied to explicit evaluation designs.

References

- Alfieri, L., Burek, P., Dutra, E., Krzeminski, B., Muraro, D., Thielen, J., and Pappenberger, F. (2013). GloFAS – global ensemble streamflow forecasting and flood early warning. *Hydrology and Earth System Sciences*, 17:1161–1175.
- Bai, Y.-T., Jia, W., Jin, X.-B., Su, T.-L., Kong, J.-L., and Shi, Z.-G. (2023). Nonstationary time series prediction based on deep echo state network tuned by bayesian optimization. *Mathematics*, 11(6):1503.
- Barata, R., Prado, R., and Sansó, B. (2022). Fast inference for time-varying quantiles via flexible dynamic models with application to the characterization of atmospheric rivers. *The Annals of Applied Statistics*, 16(1):247 – 271.
- Barata, R., Prado, R., Sansó, B., and De Leon, A. (2026). *exdqlm: Extended Dynamic Quantile Linear Models*. R package version 1.1.0, <https://CRAN.R-project.org/package=exdqlm>. doi:10.32614/CRAN.package.exdqlm.
- Barlow, R. E. and Brunk, H. D. (1972). The isotonic regression problem and its dual. *Journal of the American Statistical Association*, 67(337):140–147.

- Bonas, M., Wikle, C. K., and Castruccio, S. (2024). Calibrated forecasts of quasi-periodic climate processes with deep echo state networks and penalized quantile regression. *Environmetrics*, 35(1):e2833.
- Bondell, H. D., Reich, B. J., and Wang, H. (2010). Noncrossing quantile regression curve estimation. *Biometrika*, 97(4):825–838.
- Cardoso, F. C., Berri, R. A., Borges, E. N., Dalmazo, B. L., Lucca, G., and Dias de Mattos, V. L. (2024). Echo state network and classical statistical techniques for time series forecasting: A review. *Knowledge-Based Systems*, 293:111639.
- Carvalho, C. M., Polson, N. G., and Scott, J. G. (2010). The horseshoe estimator for sparse signals. *Biometrika*, 97(2):465–480.
- Chernozhukov, V., Fernández-Val, I., and Galichon, A. (2010). Quantile and probability curves without crossing. *Econometrica*, 78(3):1093–1125.
- Copernicus Emergency Management Service (2026). Global flood awareness system documentation. <https://confluence.ecmwf.int/display/CEMS/Global+Flood+Awareness+System>. Accessed 11 May 2026.
- Ehlers, P. J., Nurdin, H. I., and Soh, D. (2025). Improving the performance of echo state networks through state feedback.
- European Centre for Medium-Range Weather Forecasts (2026). Forecast user guide: Section 5.1 generation of the ensemble. <https://confluence.ecmwf.int/display/FUG/Section+5.1+Generation+of+the+Ensemble>. Accessed 11 May 2026.
- Fan, H.-T., Wang, W., and Jin, Z. (2017). Performance optimization of echo state networks through principal neuron reinforcement. In *2017 International Joint Conference on Neural Networks (IJCNN)*, pages 1717–1723.
- Gallicchio, C., Micheli, A., and Pedrelli, L. (2018). Design of deep echo state networks. *Neural Networks*, 108:33–47.
- Gandhi, M., Lee, K., Pan, Y., and Theodorou, E. (2018). Propagating uncertainty through the tanh function with application to reservoir computing.
- Gneiting, T. (2011). Quantiles as optimal point forecasts. *International Journal of Forecasting*, 27(2):197–207.
- Gneiting, T. and Katzfuss, M. (2014). Probabilistic forecasting. *Annual Review of Statistics and Its Application*, 1:125–151.
- Gneiting, T. and Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378.

- Gonçalves, K. C. M., Migon, H. S., and Bastos, L. S. (2020). Dynamic quantile linear models: A Bayesian approach. *Bayesian Analysis*, 15(2):335–362.
- Guerra, M., Scardapane, S., and Bianchi, F. M. (2023). Probabilistic load forecasting with reservoir computing. *IEEE Access*, 11:145989–146002.
- Harrigan, S., Zsótér, E., Cloke, H., Salamon, P., and Prudhomme, C. (2023). Daily ensemble river discharge reforecasts and real-time forecasts from the operational global flood awareness system. *Hydrology and Earth System Sciences*, 27:1–19.
- Hirpa, F. A., Salamon, P., Beck, H. E., Lorini, V., Alfieri, L., Zsótér, E., and Dadson, S. J. (2018). Calibration of the global flood awareness system (GloFAS) using daily streamflow data. *Journal of Hydrology*, 566:595–606.
- Jaeger, H. (2001). The “echo state” approach to analysing and training recurrent neural networks. Technical Report 148, GMD – German National Research Center for Information Technology, Sankt Augustin, Germany.
- Jaeger, H., Lukoševičius, M., Popovici, D., and Siewert, U. (2007). Optimization and applications of echo state networks with leaky-integrator neurons. *Neural Networks*, 20(3):335–352.
- Ji, F., Lee, J., and Rabe-Hesketh, S. (2025). Valid standard errors for Bayesian quantile regression with clustered and independent data. *Journal of Educational and Behavioral Statistics*. OnlineFirst.
- Jiang, L., Bondell, H. D., and Wang, H. J. (2014). Interquantile shrinkage and variable selection in quantile regression. *Computational Statistics & Data Analysis*, 69:208–219.
- Kim, T. and King, B. R. (2020). Time series prediction using deep echo state networks. *Neural Computing and Applications*, 32(23):17769–17787.
- Koenker, R. (2005). *Quantile Regression*. Cambridge University Press, Cambridge.
- Koenker, R. and Bassett, Jr., G. (1978). Regression quantiles. *Econometrica*, 46(1):33–50.
- Kohns, D. and Szendrei, T. (2025). Joint quantile shrinkage: A state-space approach toward non-crossing bayesian quantile models.
- Kozumi, H. and Kobayashi, G. (2011). Gibbs sampling methods for Bayesian quantile regression. *Journal of Statistical Computation and Simulation*, 81(11):1565–1578.
- Lukoševičius, M. (2012). A practical guide to applying echo state networks. In Montavon, G., Orr, G. B., and Müller, K.-R., editors, *Neural Networks: Tricks of the Trade*, pages 659–686. Springer Berlin Heidelberg, Berlin, Heidelberg, second edition.
- Lukoševičius, M. and Jaeger, H. (2009). Reservoir computing approaches to recurrent neural network training. *Computer Science Review*, 3(3):127–149.

- Lv, Z., Zhao, J., Liu, Y., and Wang, W. (2016). Use of a quantile regression based echo state network ensemble for construction of prediction intervals of gas flow in a blast furnace. *Control Engineering Practice*, 46:94–104.
- Maat, J. R., Gianniotis, N., and Protopapas, P. (2018). Efficient optimization of echo state networks for time series datasets. In *2018 International Joint Conference on Neural Networks (IJCNN)*, pages 1–7. IEEE.
- Makalic, E. and Schmidt, D. F. (2016). A simple sampler for the horseshoe estimator. *IEEE Signal Processing Letters*, 23(1):179–182.
- McDermott, P. L. and Wikle, C. K. (2019). Deep echo state networks with uncertainty quantification for spatio-temporal forecasting. *Environmetrics*, 30(3):e2553.
- Nishimura, A. and Suchard, M. A. (2023). Shrinkage with shrunken shoulders: Gibbs sampling shrinkage model posteriors with guaranteed convergence rates. *Bayesian Analysis*, 18(4):1163–1194.
- Piironen, J. and Vehtari, A. (2017). Sparsity information and regularization in the horseshoe and other shrinkage priors. *Electronic Journal of Statistics*, 11(2):5018–5051.
- Saadat, J., Farshad, M., Eliasi, H., and Shokoohi Mehr, K. (2025). Establishing echo state network in order to be used in online application. *Operations Research Forum*, 6(115).
- Singh, P. and Raman, B. (2025). Echo state networks as state-space models: A systems perspective.
- Sriram, K., Ramamoorthi, R. V., and Ghosh, P. (2013). Posterior consistency of Bayesian quantile regression based on the misspecified asymmetric laplace density. *Bayesian Analysis*, 8(2):479–504.
- Sun, C., Song, M., Cai, D., Zhang, B., Hong, S., and Li, H. (2024). A systematic review of echo state networks from design to application. *IEEE Transactions on Artificial Intelligence*, 5(1):23–37.
- Szendrei, T., Bhattacharjee, A., and Schaffer, M. E. (2025). Fused LASSO as non-crossing quantile regression.
- Viehweg, J., Teutsch, P., and Mäder, P. (2025). A systematic study of echo state networks topologies for chaotic time series prediction. *Neurocomputing*, 618:129032.
- Wang, C. and Blei, D. M. (2013). Variational inference in non-conjugate models. *Journal of Machine Learning Research*, 14:1005–1031.
- Wang, J. and Cai, Z. (2024). A composite bayesian approach for quantile curve fitting with non-crossing constraints. *Communications in Statistics – Theory and Methods*.
- Wu, Q., Fokoué, E., and Kudithipudi, D. (2018). On the statistical challenges of echo state networks and some potential remedies.

- Wu, T. and Narisetty, N. N. (2021). Bayesian multiple quantile regression for linear models using a score likelihood. *Bayesian Analysis*, 16(4):1261–1284.
- Wu, Y. and Liu, Y. (2009). Stepwise multiple quantile regression estimation using non-crossing constraints. *Statistics and Its Interface*, 2(3):299–310.
- Yan, M., Huang, C., Bienstman, P., Tiño, P., Lin, W., and Sun, J. (2024). Emerging opportunities and challenges for the future of reservoir computing. *Nature Communications*, 15:2056.
- Yan, Y., Zheng, X., and Kottas, A. (2025). A new family of error distributions for Bayesian quantile regression. *Bayesian Analysis*, pages 1–29. Advance Publication.
- Yang, Y. and Tokdar, S. T. (2017). Joint estimation of quantile planes over arbitrary predictor spaces. *Journal of the American Statistical Association*, 112(519):1107–1120.
- Yang, Y., Wang, H. J., and He, X. (2016). Posterior inference in Bayesian quantile regression with asymmetric laplace likelihood. *International Statistical Review*, 84(3):327–344.
- Yildiz, I. B., Jaeger, H., and Kiebel, S. J. (2012). Re-visiting the echo state property. *Neural Networks*, 35:1–9.
- Yu, K. and Moyeed, R. A. (2001). Bayesian quantile regression. *Statistics & Probability Letters*, 54(4):437–447.
- Yu, R., Gu, C., Stiasny, J., Wen, Q., Dilov, W. S., Qi, L., and Cremer, J. L. (2026). PriceFM: Foundation model for probabilistic electricity price forecasting. Version 4, revised 8 May 2026.
- Zhang, H. and Vargas, D. V. (2023). A survey on reservoir computing and its interdisciplinary applications beyond traditional machine learning. *IEEE Access*, 11:81033–81070.
- Zou, H. and Yuan, M. (2008). Composite quantile regression and the oracle model selection theory. *The Annals of Statistics*, 36(3):1108–1126.