

Bayesian Quantile Deep Echo State Networks for Nonlinear Time Series

Antonio De Leon¹ Raquel Prado¹ Bruno Sansó¹

¹Department of Statistics, University of California, Santa Cruz

August 28, 2026

Abstract

Conditional quantiles are central to asymmetric decision losses, tail-risk assessment, and interval forecasts, but Bayesian quantile regression for nonlinear time series can be difficult when the temporal feature vector is high-dimensional. We develop the quantile deep echo state network (Q-DESN), a Bayesian quantile regression model conditional on fixed features generated by a deep echo state network. Conditional on a specified reservoir construction and feature map, posterior uncertainty is assigned to the regression coefficients, likelihood, and shrinkage parameters. Single-level fits use asymmetric Laplace or quantile-fixed generalized asymmetric Laplace working likelihoods with ridge or regularized-horseshoe priors. Posterior computation uses Markov chain Monte Carlo when computationally practical and a model-specific variational Bayes approximation for analyses requiring repeated fitting. For quantile grids, we compare independent level-wise regressions followed by monotone rearrangement with a joint quantile-vector regression that shrinks adjacent quantile-specific coefficient differences. Synthetic studies, a single-origin retrospective Global Flood Awareness System (GloFAS) streamflow case, and a retrospective PriceFM comparison identify settings where this fixed-feature Bayesian regression improves finite-grid quantile scores.

1 Introduction

Conditional quantiles are natural summaries for asymmetric decision losses, tail-risk assessment, and interval forecasts. If $\mathcal{F}_t$ is the information available at forecast origin t , the level- p quantile of a future response is

$$Q_{t,h}(p) = \inf\{u : \Pr(Y_{t+h} \leq u \mid \mathcal{F}_t) \geq p\}, \quad 0 < p < 1.$$

Quantile regression estimates this quantity directly (Koenker and Bassett, 1978; Koenker, 2005). Dynamic applications are more difficult when nonlinear temporal structure must be represented by a large number of correlated predictors and inference must be repeated over forecast origins or probability levels.

Echo state networks (ESNs) represent temporal dependence through recurrent nonlinear states whose reservoir weights are drawn once rather than estimated with the regression coefficients (Jaeger, 2001; Lukoševičius and Jaeger, 2009). Deep ESNs stack these reservoirs to obtain features at several time scales (Gallicchio et al., 2018). We use the resulting states as a fixed nonlinear design and place Bayesian quantile regression on the final linear predictor. This separation retains the computational economy of reservoir methods while providing regularized posterior inference for the conditional quantile.

The proposed quantile deep echo state network (Q-DESN) uses either the asymmetric Laplace (AL) likelihood or the quantile-fixed generalized asymmetric Laplace likelihood of Yan et al. (2025), abbreviated exAL. Both are working likelihoods: they define posterior inference for quantile-regression parameters without asserting the corresponding error law (Sriram et al., 2013; Yang et al., 2016; Ji et al., 2025). The exAL family adds a shape parameter while preserving the chosen conditional quantile as its location. Ridge regularization supplies a dense baseline, and the regularized horseshoe (RHS) supplies adaptive global-local shrinkage with a finite slab (Carvalho et al., 2010; Piironen and Vehtari, 2017).

Several ESN methods already address probabilistic or quantile forecasting (McDermott and Wikle, 2019; Lv et al., 2016; Bonas et al., 2024). Q-DESN differs by treating a specified DESN realization as a fixed feature map and assigning a posterior to the quantile-specific regression, likelihood, and shrinkage parameters. Its linear dynamic comparators are the DQLM and exDQLM (Gonçalves et al., 2020; Barata et al., 2022), implemented in version 1.1.1 of the CRAN package `exdqmlm` (Barata et al., 2026).

For a grid of probability levels, we distinguish two inferential strategies. Independent regressions may be rearranged after fitting to obtain an ordered reported grid (Barlow and Brunk, 1972; Chernozhukov et al., 2010). A joint quantile-vector regression instead shares information during estimation by placing RHS priors on adjacent quantile-specific coefficient differences, adapting the Quantile-Varying Parameter construction of Kohns and Szendrei (2025). This prior can reduce crossings but does not enforce noncrossing; an ordered parameterization or explicit rearrangement is still required for exact monotonicity (Wu and Liu, 2009; Bondell et al., 2010; Yang and Tokdar, 2017).

The article contributes the fixed-feature Bayesian Q-DESN formulation, its independent and joint quantile-grid variants, and posterior computation by Markov chain Monte Carlo (MCMC) and a model-specific variational Bayes approximation. A single-quantile simulation compares Q-DESN with DQLM and exDQLM; a second simulation isolates joint versus independent estimation of a seven-level quantile grid. Retrospective GloFAS streamflow and PriceFM electricity-price examples assess independent Q-DESN regressions under their stated information sets. The supplement contains the complete posterior derivations, initialization and forecast details, diagnostics, and full numerical results.

2 Fixed-Feature Bayesian Quantile Models

2.1 Deep echo state network features

Let $\mathbf{u}_t = (y_{t-1}, \dots, y_{t-m}, \mathbf{z}_t^\top)^\top$ contain response lags and exogenous variables. For layer d , $\bar{\mathbf{W}}_d$ is a fixed sparse recurrent matrix scaled to target spectral radius ρ_d , $\mathbf{W}_d^{\text{in}}$ is a fixed input matrix, α_d is the leak rate, and $\mathbf{Q}_{d-1}$ reduces the preceding layer. With $\tilde{\mathbf{h}}_{t,d-1} = \mathbf{Q}_{d-1}\mathbf{h}_{t,d-1}$, the state recursion is

$$\begin{aligned} \text{First layer: } \mathbf{h}_{t,1} &= (1 - \alpha_1)\mathbf{h}_{t-1,1} + \alpha_1 f(\bar{\mathbf{W}}_1\mathbf{h}_{t-1,1} + \mathbf{W}_1^{\text{in}}\mathbf{u}_t), \\ \text{Higher layers: } \mathbf{h}_{t,d} &= (1 - \alpha_d)\mathbf{h}_{t-1,d} + \alpha_d f(\bar{\mathbf{W}}_d\mathbf{h}_{t-1,d} + \mathbf{W}_d^{\text{in}}\tilde{\mathbf{h}}_{t,d-1}), \quad d = 2, \dots, D, \quad (2.1) \\ \text{Regression features: } \tilde{\mathbf{x}}_t &= [\mathbf{h}_{t,D}^\top, k(\tilde{\mathbf{h}}_{t,1})^\top, \dots, k(\tilde{\mathbf{h}}_{t,D-1})^\top, \mathbf{u}_t^\top]^\top, \quad \mathbf{x}_t = (1, \tilde{\mathbf{x}}_t^\top)^\top. \end{aligned}$$

The reservoir quantities in Eq. (2.1), including scaling and washout, are determined from the fitting information and then held fixed. Posterior inference is therefore conditional on $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_T)^\top$. The supplement gives the random weight construction, dimensions, initialization, and washout conventions.

2.2 Single-level Q-DESN

At probability level p_0 , the marginal working model is

$$y_t \mid \mathbf{x}_t, \beta, \sigma, \gamma \sim \text{exAL}_{p_0}(\mathbf{x}_t^\top \beta, \sigma, \gamma), \quad t = 1, \dots, T, \quad (2.2)$$

where $\mathbf{x}_t^\top \beta$ is the fitted conditional p_0 -quantile, $\sigma > 0$ is a scale, and $\gamma \in (L, U)$ controls shape. The quantile-fixing construction determines (L, U) and functions $A_{p_0}(\gamma)$, $B_{p_0}(\gamma)$, and $D_{p_0}(\gamma)$, where the latter is the exAL positive-shift coefficient. Computation uses

$$\begin{aligned} \text{Latent variables: } v_t \mid \sigma &\sim \text{Exp}(\text{rate} = 1/\sigma), \quad s_t \sim \mathcal{N}^+(0, 1), \\ \text{Gaussian conditional: } y_t \mid \cdot &\sim \mathcal{N}\left(\mathbf{x}_t^\top \beta + D_{p_0}(\gamma)\sigma s_t + A_{p_0}(\gamma)v_t, B_{p_0}(\gamma)\sigma v_t\right). \end{aligned} \quad (2.3)$$

Integrating out (s_t, v_t) recovers Eq. (2.2). The AL model is the nested case with its probability-specific constants and no s_t or γ block. Posterior intervals and forecast bands are thus model-based summaries conditional on the fixed feature map and working likelihood; their empirical coverage is assessed separately.

The intercept has a weak Gaussian prior. For the r non-intercept coefficients, ridge regularization uses $\beta_{\text{slope}} \sim \mathcal{N}(\mathbf{0}_r, \kappa_\beta^2 \mathbf{I}_r)$. The adaptive alternative uses the Nishimura–Suchard product form of the regularized horseshoe (Nishimura and Suchard, 2023; Makalic and Schmidt, 2016). Conditional on global, local, and slab scales (τ, λ_j, ζ) ,

$$\beta_j \mid \tau, \lambda_j, \zeta \sim \mathcal{N}(0, V_j), \quad V_j = \left(\zeta^{-2} + \tau^{-2}\lambda_j^{-2}\right)^{-1}, \quad (2.4)$$

with $\lambda_j \sim C^+(0, 1)$ and $\tau \sim C^+(0, \tau_0)$. The product density and all scale-dependent normalizing factors are stated in the supplement; Eq. (2.4) also shows how the slab bounds the effective coefficient variance.

Reference calibration of the global scale. The hyperprior scale τ_0 controls overall shrinkage. With standardized predictors, N_{fit} fitting observations, an elicited reference $m_0 \in (0, r)$ for the number of coefficients expected to escape strong shrinkage, and residual scale σ_{ref} , the Gaussian reference calibration of Piironen and Vehtari (2017) is

$$\tau_{0,\text{ref}} = \frac{m_0}{r - m_0} \frac{\sigma_{\text{ref}}}{\sqrt{N_{\text{fit}}}}. \quad (2.5)$$

We recommend the scale in Eq. (2.5) when comparing the Gaussian, AL, and exAL models on the same standardized design; it fixes the coefficient-prior scale while each likelihood retains its own posterior. Likelihood-information calibrations and separate choices for joint-model coefficient blocks are given in the supplement.

2.3 Joint quantile-vector regression

For $0 < p_1 < \dots < p_K < 1$, the joint model uses

$$Q_{p_k}(y_t \mid \mathcal{F}_t) = \beta_{0k} + \tilde{\mathbf{x}}_t^\top \boldsymbol{\beta}_k, \quad k = 1, \dots, K. \quad (2.6)$$

The joint regression in Eq. (2.6) multiplies the K quantile-specific AL or exAL contributions. This is an untempered composite working likelihood for repeated use of the response series. Let $\boldsymbol{\beta}_{\text{base}}$ be a baseline slope and define

$$\boldsymbol{\Delta}_1 = \boldsymbol{\beta}_1 - \boldsymbol{\beta}_{\text{base}}, \quad \boldsymbol{\Delta}_k = \boldsymbol{\beta}_k - \boldsymbol{\beta}_{k-1}, \quad k = 2, \dots, K.$$

Separate RHS hierarchies are assigned to the baseline and to each $\boldsymbol{\Delta}_k$. This distinguishes effects shared across levels from departures between neighboring levels. For $K = 1$, the baseline block is omitted and the RHS prior is placed directly on $\boldsymbol{\Delta}_1 = \boldsymbol{\beta}_1$.

Adjacent-difference shrinkage favors similar neighboring slopes, but the gap

$$Q_{p_k}(y_t \mid \mathcal{F}_t) - Q_{p_{k-1}}(y_t \mid \mathcal{F}_t) = (\beta_{0k} - \beta_{0,k-1}) + \tilde{\mathbf{x}}_t^\top (\boldsymbol{\beta}_k - \boldsymbol{\beta}_{k-1})$$

need not be positive. We consequently report crossings before rearrangement and apply an explicit monotone reporting rule when ordered output is required. The stacked Gaussian representation, block precision, and full joint posterior are given in the supplement.

3 Posterior Computation

The auxiliary representation in Eq. (2.3) yields Gaussian coefficient updates, generalized inverse-Gaussian mixture updates, and inverse-gamma RHS scale updates. For exAL, the scale–asymmetry block remains nonconjugate. MCMC updates it on transformed coordinates; the independent simulation uses a partially collapsed transition that integrates out the scale while updating γ and then redraws the scale from its full conditional. The joint model applies the corresponding update separately at each probability level. Convergence assessment includes rank-normalized $\hat{R}$, effective sample sizes, chain comparisons of reported functionals, and quantile-crossing diagnostics.

For analyses requiring many repeated fits, we use a mean-field variational approximation to the same augmented posterior. Gaussian, GIG, truncated-Normal, and inverse-gamma factors retain closed-form coordinate updates. The nonconjugate exAL scale–asymmetry factor uses the Laplace–Delta approximation of Wang and Blei (2013); we denote this model-specific procedure by VB–LD. The evidence lower bound, full conditional distributions, algorithms, and joint sparse Gaussian calculation are collected in the supplement. Variational intervals are approximate and are identified separately from MCMC posterior intervals.

4 Forecasting, Scoring, and Model Selection

At forecast origin T , posterior or variational draws are propagated through the fixed reservoir recursion using covariates available at the origin or specified by the analysis. If future response lags are required, they are generated under the stated recursion or supplied by a stated scenario. For a single-level draw, the fitted quantile location at horizon h is $\mu_{p,T+h}^{(b)} = \mathbf{x}_{T+h}^{(b)\top} \boldsymbol{\beta}_p^{(b)}$. This is a draw of the fitted conditional quantile. Sampling the AL or exAL working likelihood instead produces a response trajectory; the two summaries are not interchangeable. The supplement gives the complete recursive constructions.

For independent fits at $0 < p_1 < \dots < p_K < 1$, with $K \geq 2$, the raw estimates may cross. We write $q_{T,h,k}^*$ for the reported value after a monotone transformation when one is applied. The simulation analyses use equal-weight isotonic projection,

$$\mathbf{q}_{T,h}^* = \arg \min_{z_1 \leq \dots \leq z_K} \sum_{k=1}^K (\hat{q}_{T,h,k} - z_k)^2. \quad (4.1)$$

The projection in Eq. (4.1) is applied at the study-specific choice of draw level or summary level stated with the results. In analyses based on summary-level projection, the posterior-mean quantile curve or variational posterior-mean quantile estimates are transformed before scoring. Linear interpolation defines the quantile curve only over $[p_1, p_K]$; a complete predictive distribution also requires a tail rule.

Let the check loss be

$$\rho_p(e) = e\{p - \mathbf{1}(e < 0)\}. \quad (4.2)$$

For a predictive distribution F with finite first moment, the CRPS has the integrated check-loss representation (Gneiting and Raftery, 2007)

$$\text{CRPS}(F, y) = 2 \int_0^1 \rho_p\{y - F^{-1}(p)\} dp. \quad (4.3)$$

For the finite evaluation grid $p_1, \dots, p_K$, we approximate Eq. (4.3) by the trapezoidal rule over its covered range,

$$\begin{aligned} \text{Finite-grid score: } \text{aCRPS}_{p_1:p_K}(y; \mathbf{q}^*) &= 2 \sum_{k=1}^K \omega_k \rho_{p_k}(y - q_k^*), \\ \text{Trapezoidal weights: } \omega_1 &= \frac{p_2 - p_1}{2}, \quad \omega_k = \frac{p_{k+1} - p_{k-1}}{2}, \quad 2 \leq k \leq K-1, \\ \omega_K &= \frac{p_K - p_{K-1}}{2}. \end{aligned} \quad (4.4)$$

The weights satisfy $\sum_{k=1}^K \omega_k = p_K - p_1$. Thus aCRPS retains the scale of integrated check loss over the fitted range rather than representing a full-distribution CRPS.

In the joint simulation the true conditional response distribution is known. For N_{eval} held-out origin-horizon pairs, posterior draw b , and rearranged quantiles $q_{\ell,k}^{*,(b)}$, we also use

$$S_{\text{DGP}}^{(b)} = \frac{2}{N_{\text{eval}}} \sum_{\ell=1}^{N_{\text{eval}}} \sum_{k=1}^K \omega_k \mathbb{E}_{Y_\ell \sim F_{0,\ell}} \left[\rho_{p_k}\{Y_\ell - q_{\ell,k}^{*,(b)}\} \right]. \quad (4.5)$$

This finite-grid score averages over the known conditional response distribution at the fixed held-out design points. Its interval excludes repeated-simulation and model-selection uncertainty; the supplement records the coupling used for independently fitted levels.

For a statistical model M , reservoir candidates $R_1, \dots, R_J$ are compared by a stated validation criterion,

$$\hat{R}_M \in \arg \min_{R_j: j=1, \dots, J} S_{\text{val}}(R_j; M). \quad (4.6)$$

After selection by Eq. (4.6), the specification is fixed before held-out evaluation. The response convention, information set, forecast origins and horizons, transformations, feature ordering, and score remain common within a comparison. State checks cover finite values, activation saturation, near-constant or highly dependent columns, effective rank, conditioning, and variation across reservoir realizations. These checks are needed because spectral-radius scaling alone does not guarantee suitable driven dynamics (Jaeger et al., 2007; Yildiz et al., 2012).

The joint simulation and both applications use sequential initialization centered at the median. A Gaussian ridge fit initializes Gaussian RHS-VB; compatible estimates then initialize the median AL regression. AL fits proceed from the median toward each tail, each exAL fit starts from the matching AL level, and the complete independent grids initialize their corresponding joint models. A VB fit initializes MCMC only for the same statistical model. These steps alter starting values, not the destination likelihood, prior, posterior, or independence of level-wise fits. The independent

validation study uses a different procedure. The complete sequence, compatibility conditions, and diagram are provided in the supplement.

5 Simulation Studies

The simulation study addresses two complementary questions. The first asks whether a single-quantile Q-DESN regression can recover and forecast a known dynamic quantile path under a 500-observation benchmark following the family-based comparison strategy of Yan et al. (2025). The target quantile level is fixed in advance, the data-generating mechanism is known, and competing models are compared under matched error families, fitting-sample sizes, and quantile levels. The second asks whether a joint multi-quantile regression can recover several known conditional quantile paths under posterior simulation and reduce adjacent-level crossings before monotone rearrangement. The supplement reports the corresponding VB analyses and additional forecast comparisons.

In contrast to the linear-regression setting of Yan et al. (2025), the true conditional quantile paths, referred to as oracle paths in this simulation, are dynamic. Q-DESN represents those paths through high-dimensional deterministic reservoir features. The design therefore separates two questions: how accurately a fitted model recovers the known conditional quantile path, and how well it forecasts held-out dynamic blocks. We use p for a single target level and p_k , $k = 1, \dots, K$, for a multi-quantile grid.

The main-text tables report MCMC summaries under fixed evaluation designs: data-generating mechanisms, rolling-origin definitions, fitting samples, held-out periods, scoring rules, and model classes are specified before the comparisons are formed. The supplement gives corresponding approximate VB summaries and additional diagnostic summaries for the same simulation settings. The single-quantile and multi-quantile studies use separate fixed designs. Their results address distinct comparisons and provide complementary evidence.

5.1 Single-Quantile Dynamic Data-Generating Processes

For each error family f and target quantile level $p \in \{0.05, 0.25, 0.50\}$, the synthetic series and its conditional quantile path are

$$\begin{aligned} \text{Centered innovation: } \varepsilon_{t,p,f} &= e_{t,f} - H_f^{-1}(p), \\ \text{Observation: } y_{t,p,f} &= \mu_{t,f} + \varepsilon_{t,p,f}, \\ \text{Conditional quantile: } Q_p(y_{t,p,f} \mid \boldsymbol{\theta}_{t,f}) &= \mu_{t,f} = \mathbf{F}^\top \boldsymbol{\theta}_{t,f}, \quad \mathbf{F} = (1, 0, 1, 0, 1, 0)^\top. \end{aligned} \tag{5.1}$$

In Eq. (5.1), H_f is the distribution function of the raw error $e_{t,f}$, so the centering makes the p -quantile of $\varepsilon_{t,p,f}$ zero. The six-dimensional state contains a local level, local slope, and two sine-cosine harmonic pairs. The seasonal period is 90, the harmonics are 1 and 2, and the state innovation covariance satisfies $\sqrt{\text{diag}(\mathbf{W}_\theta)} = (0.005, 0.00002, 0.004, 0.004, 0.003, 0.003)$. The initial covariance is $\mathbf{C}_0 = 0.01\mathbf{I}_6$, and each simulated trajectory is initialized deterministically at the

family-specific state in Table 1. For a fixed family, the same latent path $\{\mu_{t,f}\}$ is used across the three values of p ; only the centering constant $H_f^{-1}(p)$ changes.

Family	Raw error distribution H_f	Initial level	Initial slope	Seasonal settings
Gaussian	$\mathcal{N}(0, 10^2)$	40	0.012	harmonics (24, 0.35), (8, -0.80)
Laplace	Laplace with scale 10	35	0.011	harmonics (28, -0.15), (10, 0.75)
Gaussian mixture	$0.1\mathcal{N}(0, 0.5^2) + 0.9\mathcal{N}(1, 15^2)$	45	0.014	harmonics (32, 0.85), (12, -1.25)

Table 1: Single-quantile dynamic data-generating mechanisms. Raw errors are shifted by $H_f^{-1}(p)$ so that the p -quantile of the observation error is zero; the same latent dynamic quantile path is then shared across target levels within each family. Seasonal entries report amplitude and phase for the two harmonic pairs in the initial state.

5.2 Single-Quantile Fit and Forecast Evaluation

Each synthetic series is first run for $T_{\text{warm}} = 2000$ warmup observations, which are discarded before analysis. The next $T_{\text{main}} = 10000$ observations form the retained series. The first 9000 support model specification and fitting, and the final 1000 are held out for forecast evaluation. The reported single-quantile comparison uses $T_{\text{fit}} = 500$, namely observations 8501–9000 immediately preceding the fitting origin. Forecast scores are computed from rolling origins within the held-out period; origins are spaced 30 observations apart, and each origin is scored for horizons up to 30 steps ahead. The same evaluation design is used for all model specifications and score calculations.

5.3 Single-Quantile Competing Methods

For each family and target quantile level in the single-quantile simulation comparison, the reported Q-DESN rows vary the working likelihood and inference method while using regularized-horseshoe (RHS) coefficient shrinkage. Entries labeled Q-DESN AL-RHS use the AL working likelihood, whereas entries labeled Q-DESN exAL-RHS use the exAL working likelihood. The main comparison focuses on regularized-horseshoe fits. For each displayed metric, the reservoir specification and realization were selected before held-out evaluation and then held fixed. Dynamic quantile linear model (DQLM) and extended dynamic quantile linear model (exDQLM) baselines are fit to the same family, quantile level, fitting-sample size, and forecast-evaluation design using version 1.1.1 of the CRAN package `exdqml` (Barata et al., 2026). The exDQLM calculations use its partially collapsed MCMC update and structured $q(\gamma)q(\sigma \mid \gamma)$ variational approximation for the scale-asymmetry block. For rolling-origin MCMC evaluation, the initial parameter posterior is retained while the filtering distribution is updated sequentially using the responses available at each origin. In the exDQLM analysis, the observation-error mean and variance used in this update are averaged over

paired retained draws of the scale and asymmetry parameters. The updated state is then propagated over the requested forecast horizons.

The primary comparison is a fixed-design benchmark against structured dynamic quantile linear models under linear dynamic mechanisms that favor those baselines. Repeated data-generating realizations, broader reservoir-parameter comparisons, and non-Bayesian ESN regression baselines remain topics for future sensitivity analyses.

5.4 Single-Quantile Criteria for Comparison

Let $q_{p,t}^{(b)}$ be retained draw b of the conditional p -quantile at fitting index t , and let $q_{p,\ell}^{(b)}$ be the corresponding forecast draw for origin-horizon pair ℓ . With $N_{\text{fit}} = 500$ and $N_{\text{eval}} = 1000$, the draw-wise criteria are

$$\text{Fit RMSE: } R_{\text{fit}}^{(b)} = \left\{ \frac{1}{N_{\text{fit}}} \sum_{t=1}^{N_{\text{fit}}} (q_{p,t}^{(b)} - \mu_t)^2 \right\}^{1/2}, \quad (5.2)$$

$$\text{Forecast MAE: } A_{\text{fore}}^{(b)} = \frac{1}{N_{\text{eval}}} \sum_{\ell=1}^{N_{\text{eval}}} |q_{p,\ell}^{(b)} - \mu_\ell|, \quad (5.3)$$

$$\text{Forecast check loss: } C_{\text{fore}}^{(b)} = \frac{1}{N_{\text{eval}}} \sum_{\ell=1}^{N_{\text{eval}}} \rho_p(y_\ell - q_{p,\ell}^{(b)}), \quad (5.4)$$

where ρ_p is the check loss in Eq. (4.2), and μ_ℓ and y_ℓ denote the oracle quantile and realized outcome for pair ℓ . Thus fit RMSE and forecast MAE measure recovery of the oracle quantile path, whereas forecast check loss is a proper score against realized observations.

For each criterion $M^{(b)}$, the tables report the posterior mean $B^{-1} \sum_{b=1}^B M^{(b)}$ and the empirical 0.025 and 0.975 posterior quantiles. The VB panels use draws from the variational approximation and are therefore approximate. Because the criteria are nonlinear functions of the quantile path, these posterior means generally differ from applying a criterion once to a posterior point path.

For each model, error family, quantile level, and criterion, the reservoir specification was selected in a separate calibration analysis before the posterior metric distribution was evaluated. Consequently, different criteria within a model may correspond to different reservoir specifications. The intervals condition on the selected specification, simulated series, rolling-origin design, and reservoir realization; repeated-simulation, reservoir-selection, and model-selection uncertainty are excluded.

5.5 Single-Quantile Fit and Forecast Results

Figures 1 and 2 summarize the MCMC posterior distributions of the two forecast criteria. Among the 27 family-quantile-criterion comparisons, Q-DESN exAL-RHS has the lowest posterior mean in 10 and Q-DESN AL-RHS in 8, compared with 4 for DQLM and 5 for exDQLM. The equal-tailed intervals of the two lowest means overlap in every comparison.

The displayed bands are posterior distributions of the draw-wise aggregate criteria. For Q-DESN and exQ-DESN, each metric draw preserves a common posterior draw of the quantile-regression coefficients across forecast origins and horizons. For DQLM and exDQLM, the origin- and horizon-specific posterior quantile marginals are combined using a pre-specified product coupling. The resulting intervals quantify model-specific posterior uncertainty for the reported criteria, conditional on the simulated series, selected specification, and evaluation design.

Posterior means were numerically unchanged across coupling assumptions, whereas interval widths varied; we therefore report the product-coupling intervals with this sensitivity analysis.

All preprocessing uses fitting information only. Complete MCMC and approximate VB values, fit-recovery results, and a five-chain sensitivity analysis are reported in the supplement.

5.6 Joint Multi-Quantile Study with a Shared Feature Design

The second simulation study compares joint and independent estimation at the seven levels $p_k \in \{0.05, 0.10, 0.25, 0.50, 0.75, 0.90, 0.95\}$ under the AL and exAL working likelihoods. Eight mechanisms span Gaussian, Laplace, Gaussian-mixture, asymmetric-tail, persistent-heavy-tail, Student- t , regime-shift, and nonlinear dynamics. For each mechanism, one DESN and RHS specification is selected using separate source realizations and then shared by all four model classes. The 500-observation evaluation realization is not used for this selection. Thus differences among the 32 comparisons are not confounded by model-specific feature maps or prior scales.

The evaluation fits use sequential initialization centered at the median. A Gaussian ridge fit initializes Gaussian RHS-VB; the resulting compatible estimates initialize the median AL regression; and the remaining independent AL fits proceed toward the two tails. Each exAL fit is initialized from the AL fit at the same level, the complete independent grids initialize the corresponding joint fits, and each completed VB fit initializes MCMC for the same model. These transfers determine starting values only. Within each model cell, all five MCMC chains share fixed prior hyperparameters and the same posterior target. The supplement gives the complete calculation and chain allocation.

Held-out evaluation is sequentially conditional: regression coefficients and reservoir weights remain fixed, while newly observed response lags enter later forecast origins. Each comparison contains the same 990 origin-horizon rows. The primary criterion is S_{DGP} in Eq. (4.5), computed for each retained MCMC draw after monotone rearrangement. Independent quantile draws are combined by the pre-specified chain-balanced product coupling described in the supplement.

Figure 3 shows the corrected posterior score comparison. Independent exQDESN has the smallest numerical posterior mean in seven mechanisms; joint exQDESN is smallest for the Laplace mechanism. The canonical reported quantile grid instead favors independent exQDESN in all eight mechanisms, including the Laplace mechanism. The marginal intervals for the numerical winner and runner-up overlap in every mechanism. Of the 16 paired joint-minus-independent contrasts, four intervals favor independent estimation and twelve include zero; none favors joint estimation. The evidence therefore does not support a general forecast advantage from joint estimation under this shared design.

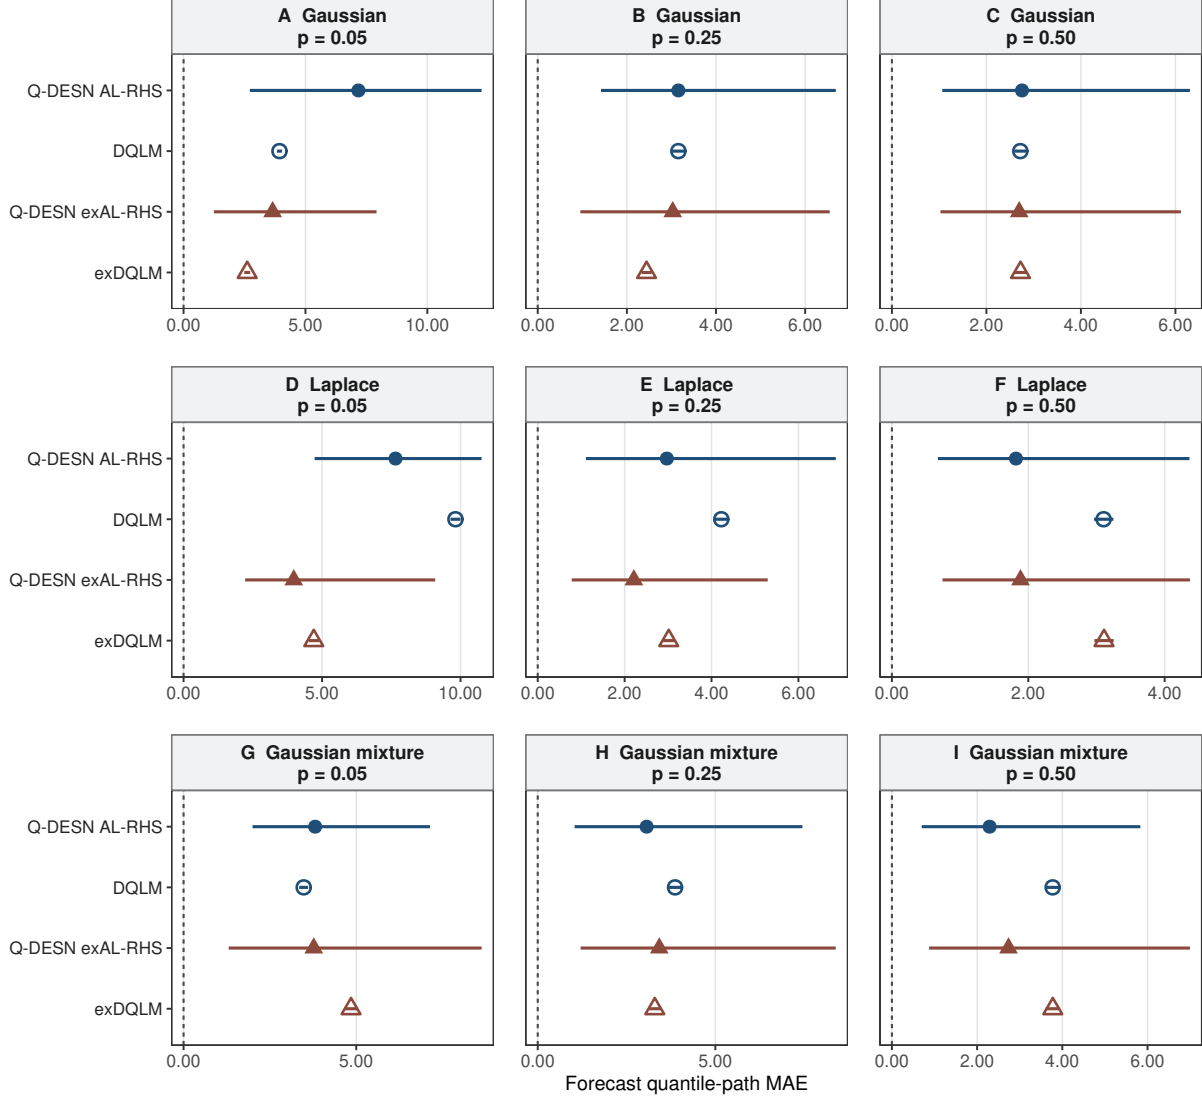


Figure 1: MCMC forecast quantile-path mean absolute error (MAE) in the single-quantile simulation. Points are posterior means and horizontal segments are equal-tailed 95% posterior intervals; lower values are better. Navy circles identify the AL models and muted-brick triangles the exAL models; filled symbols denote Q-DESN and open symbols the corresponding dynamic quantile linear model. Panels A–I identify the data-generating family and probability level and use separate horizontal scales. The dashed line marks the exact data-generating-process value of zero. Intervals condition on the simulated series, evaluation design, and case-specific selected specification.

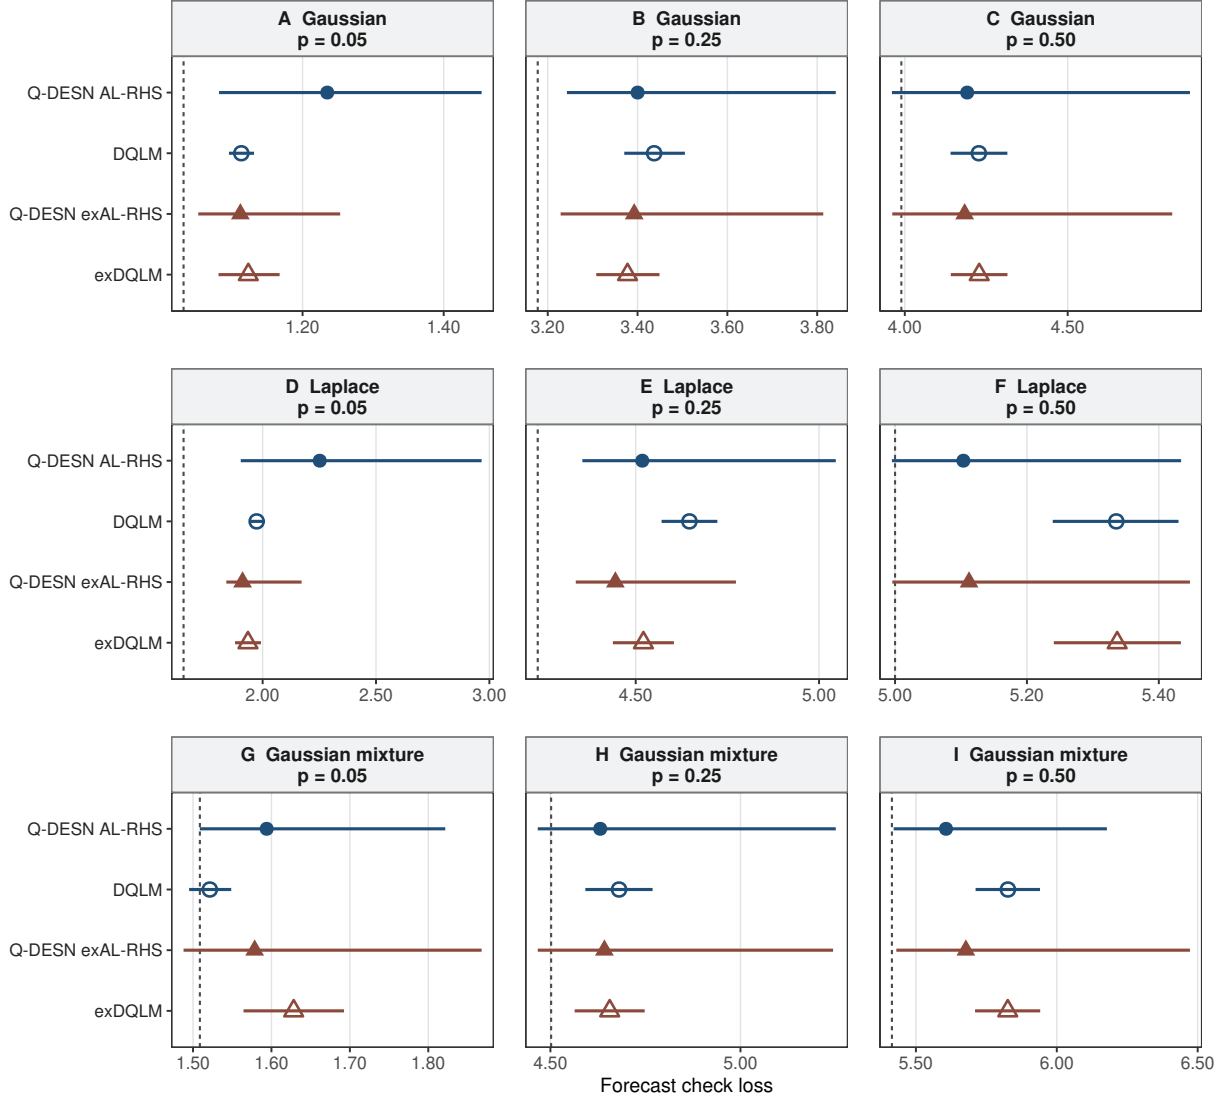


Figure 2: MCMC forecast check loss in the single-quantile simulation. Points are posterior means and horizontal segments are equal-tailed 95% posterior intervals; lower values are better. Model encodings and panels are as in Figure 1, with separate horizontal scales. The dashed line marks population expected check loss at the true conditional quantile; conditional finite-sample intervals may cross this reference.

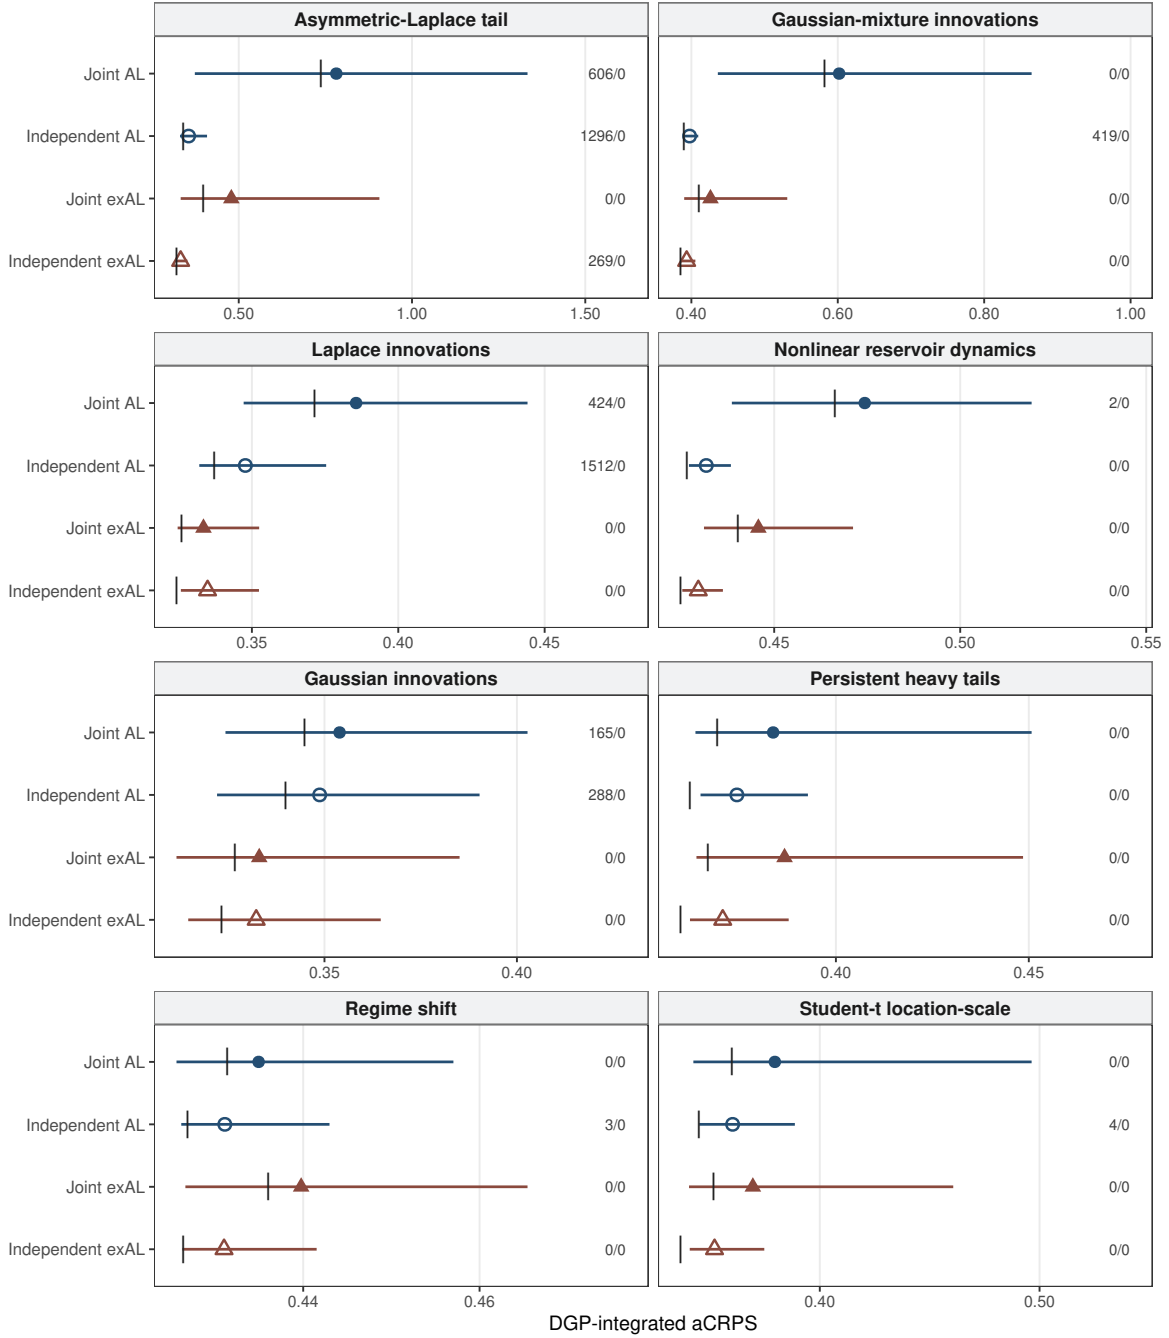


Figure 3: Corrected posterior DGP-integrated aCRPS for the joint multi-quantile study. Colored points are posterior means and horizontal segments are equal-tailed 95% intervals; black vertical marks give the canonical-action scores. Lower values are better. Filled symbols denote joint fits, open symbols independent fits, blue AL, and red exAL. Labels give raw/reported crossing counts for the canonical grid. All reported counts are zero after the pre-specified monotone rule. Marginal intervals for the numerical winner and runner-up overlap in every setting.

Before rearrangement, the canonical posterior-mean forecast grids contain 1,197 crossings for joint Q-DESN, 3,522 for independent Q-DESN, none for joint exQDESN, and 269 for independent exQDESN, summed over the eight mechanisms. Joint exQDESN is therefore the most coherent raw specification in this comparison, although its predictive score is not uniformly lowest. Rearrangement produces ordered grids in all 32 comparisons. Posterior-score diagnostics meet the pre-specified criteria in 22 comparisons and warrant review in 10; all 16 exAL scale-asymmetry assessments remain at review level. Exact scores, fit recovery, crossing rates, and computational diagnostics are reported in the supplement. A separate 50-replicate variational analysis is retained as a sensitivity study.

6 Application: GloFAS Retrospective Streamflow Quantile Forecasting

Retrospective case study. We consider one medium-range retrospective streamflow experiment using the Global Flood Awareness System (GloFAS), which is jointly developed by the European Commission and the European Centre for Medium-Range Weather Forecasts and operated within the Copernicus Emergency Management Service (Alfieri et al., 2013; Copernicus Emergency Management Service, 2026). GloFAS provides ensemble river-discharge forecasts for flood early warning and verification (Harrigan et al., 2023). The analysis estimates the conditional quantiles of a United States Geological Survey (USGS) gauge record while using retrospective and issued GloFAS values to estimate a time-varying discrepancy.

6.1 Data and Forecast-Origin Design

Let T be the forecast origin, $y_t = \log(y_t^{\text{USGS}} + 1)$ the transformed gauge record through T , g_t^{ret} the transformed retrospective GloFAS series, and $g_{T+h,e}^{\text{ens}}$ transformed issued ensemble member e at horizon h . The origin is 2022-12-25. The requested window contains 30 days, from 26 December 2022 through 24 January 2023, but the issued GloFAS record contains 28 evaluable horizons through 22 January 2023. All 51 members are retained at each available horizon. Future USGS values are excluded from fitting and used only for scoring. The forecast-period precipitation and soil-moisture covariates are realized PRISM/ERA5 values. The experiment is therefore a retrospective oracle-covariate diagnostic, not a real-time forecast based only on information available at T .

6.2 Application Model

Latent-path discrepancy model. Two fixed DESN feature maps represent the reference process and the GloFAS-minus-USGS discrepancy. For probability level p , write $q_{p,t} = \mathbf{x}_t^{\text{ref}\top} \boldsymbol{\beta}_p$ for the reference quantile and $d_{p,t} = \mathbf{x}_t^{\text{disc}\top} \boldsymbol{\delta}_p$ for the additive discrepancy. On the transformed scale, the

historical and issued working-likelihood contributions are

$$\begin{aligned}
&\text{Reference history: } y_t \mid \cdot \sim \text{AL}_p(q_{p,t}, \sigma_Y), \\
&\text{Retrospective GloFAS: } g_t^{\text{ret}} \mid \cdot \sim \text{AL}_p(q_{p,t} + d_{p,t}, \sigma_G), \\
&\text{Issued member: } g_{T+h,e}^{\text{ens}} \mid \cdot \sim \text{AL}_p(q_{p,T+h} + d_{p,T+h}, \sigma_G).
\end{aligned} \tag{6.1}$$

The unavailable path $y_{T+1:T+30}$ is represented as latent during fitting and recursively supplies future response lags to the reference reservoir. Each of the 51 issued members has weight $1/51$, so every issued horizon has total weight one. This weighting is applied to the augmented complete-data working likelihood.

The reported model estimates the seven levels 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.95 jointly by VB, with AL working likelihoods and adjacent-quantile RHS shrinkage for the reference and discrepancy coefficient paths. No crossing correction is applied; the selected issued-window grid has no crossings. The supplement gives the source weighting, reservoir specifications, latent-path approximation, and convergence record.

6.3 Retrospective Application Results

The selected Joint AL estimate completed ten outer sweeps. All inner fits and RHS release conditions were satisfied, and all 14 reference and discrepancy RHS blocks changed after release. Its final maximum parameter change was 0.004154, above the pre-specified 0.001 outer tolerance. We therefore describe it as cap-stabilized after RHS release, not as strictly outer-converged. No MCMC application fit is reported.

Forecast-period scores. Table 2 compares the selected Joint AL model with the completed independent quantile regressions, two Gaussian predictive baselines, and raw GloFAS. Joint AL reduces aCRPS from 1.4360 for raw GloFAS to 0.8374, a 41.7% reduction, and has the smallest mean check loss and interval score among the quantile-regression models. Independent AL gives similar scores and the same empirical 90% coverage of 0.929. Independent exAL has larger scores and lower coverage in this forecast window.

The Normal Ridge baseline has the smallest numerical aCRPS, 0.8062, but its empirical 90% coverage is 0.250 and its interval score is 13.8539, compared with 0.929 and 5.3889 for Joint AL. This difference illustrates that the summaries answer distinct questions: mean check loss weights the seven levels equally, whereas aCRPS uses the trapezoidal weights in Eq. (4.4), and interval score and coverage assess the central forecast interval. The comparisons are descriptive for one forecast origin; no sampling-uncertainty calculation for score differences is available.

This issued-window improvement does not extend to the historical fit. On the common six-level grid, the Part 4 Joint AL aCRPS is 68.6% larger than FR09 over the full history and is 104.2%, 51.4%, and 8.6% larger over the final 1000, 200, and 50 dates, respectively. FR09 is therefore retained as the stronger historical benchmark, while Part 4 is interpreted only for this single-origin issued-window experiment.

Figure 4 places the reported quantile paths against the observed history, held-out gauge record, and issued GloFAS ensemble over the single forecast window.

Table 2: Forecast scoring for the GloFAS retrospective analysis at origin 2022-12-25. Scores use the 28 issued horizons and are computed on the $\log(\text{USGS} + 1)$ scale and the common probability grid 0.05, 0.20, 0.35, 0.50, 0.65, 0.80, 0.95. Mean check loss averages the seven level-specific losses equally, whereas aCRPS uses the trapezoidal weights in Eq. (4.4). The Normal rows summarize Gaussian predictive distributions on the same dates and probability grid. Lower scores are better. Coverage is empirical coverage of the nominal 90% interval; the final column gives the aCRPS reduction relative to raw GloFAS.

Model	Mean check loss	Interval score	aCRPS	90% coverage	aCRPS reduction vs raw GloFAS
Joint AL Q-DESN (selected)	0.4180	5.3889	0.8374	0.929	41.7%
Independent AL Q-DESN	0.4292	5.4239	0.8606	0.929	40.1%
Independent exAL Q-DESN	0.5227	6.6540	1.0478	0.750	27.0%
Normal Ridge DESN	0.4334	13.8539	0.8062	0.250	43.9%
Normal RHS/VB DESN	0.5077	16.6741	0.9412	0.179	34.5%
Raw GloFAS	0.7713	24.5037	1.4360	0.000	—

7 Application: PriceFM Retrospective Electricity-Price Comparison

The second application compares retrospective Q-DESN day-ahead electricity-price forecasts with version 4 of PriceFM (Yu et al., 2026). The release covers European bidding regions at 15-minute resolution from 1 January 2022 through 1 January 2026 and includes price, load, solar, and wind forecasts. Average quantile loss (AQL) is evaluated on the PriceFM probability grid at matched held-out timestamps for 114 region-fold comparisons.

7.1 Comparison Design

The comparison contains 114 region-fold evaluations, spanning 38 European bidding regions and 3 rolling folds. Each evaluation is scored on the fold-aligned held-out timestamps and probability levels 0.10, 0.25, 0.45, 0.50, 0.55, 0.75, 0.90. Throughout this section, Δ AQL denotes Q-DESN AQL minus PriceFM AQL. Negative values indicate lower Q-DESN AQL; positive values indicate lower PriceFM AQL.

Retrospective covariates. We compare Q-DESN forecasts with the released PriceFM forecasts at matched region-fold timestamps. Across the comparison, 58 evaluations use own-region predictors and 56 evaluations also include neighboring-region summaries. Both predictor sets use retrospectively observed own-region load, solar, and wind lead covariates in the comparison; evaluations with neighborhood summaries add retrospectively observed neighboring-region lead summaries. The score comparison is therefore retrospective. A PriceFM study of the joint quantile-vector RHS prior would require the same fold-aligned design with that joint regression fitted directly.

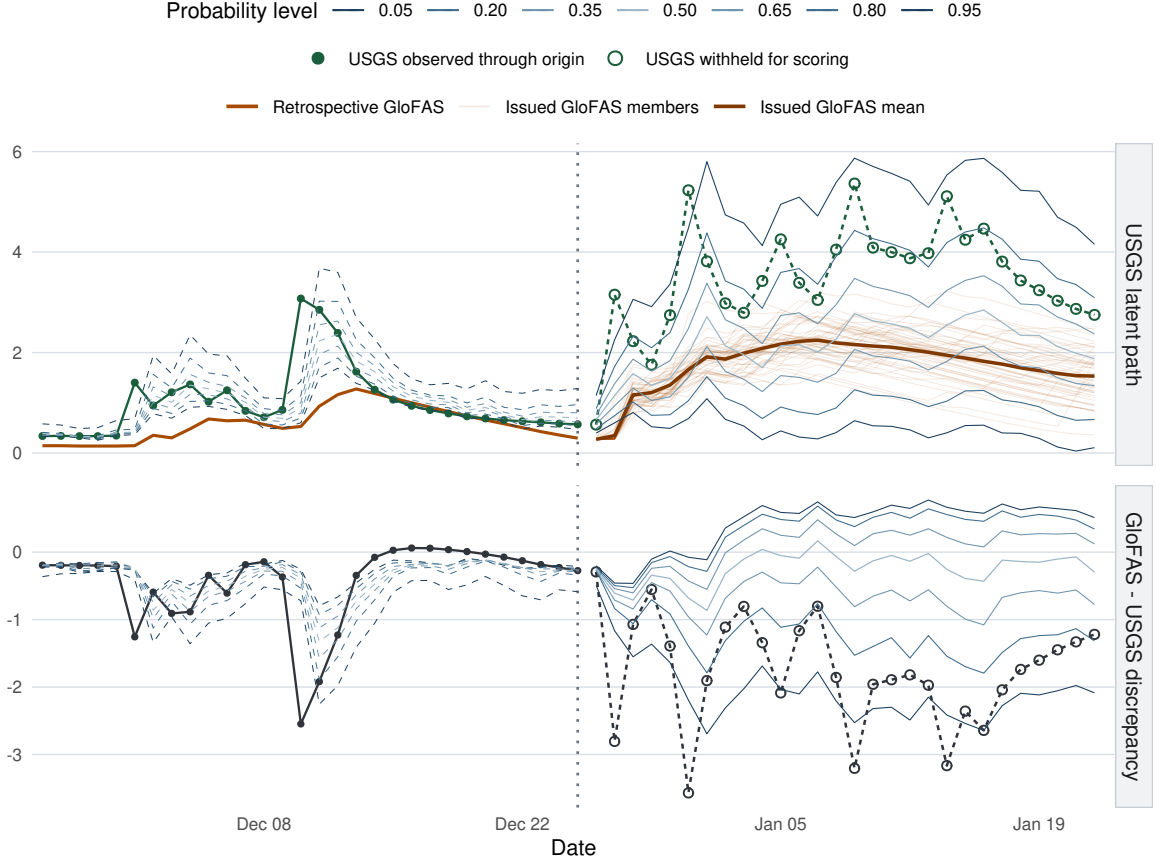


Figure 4: Joint AL Q-DESN latent-path analysis over the final 30 historical dates and 28 issued GloFAS horizons. The seven probability levels share a single-blue scale that darkens toward both tails; historical quantile paths are dashed and issued paths are solid. Forest-green filled circles joined by a thin solid line denote USGS observations available through the origin. Forest-green open circles joined by a short-dashed line denote the withheld USGS values used only for scoring. Dark-orange lines show retrospective GloFAS and the 51 issued ensemble members. The lower panel gives the corresponding GloFAS-minus-USGS discrepancy paths. The vertical dotted line marks the forecast origin. No crossing correction is applied. All values are on the transformed scale.

7.2 Results by Region, Fold, and Forecast Lead

Across the aligned benchmark, Q-DESN has lower AQL in 66 of 114 region-fold evaluations (57.9%). Relative to its own AQL, PriceFM has lower AQL by at most 5% in 12 comparisons and by more than 5% in 36 comparisons. The mean Q-DESN AQL is 6.824, compared with mean PriceFM AQL 7.039, giving mean Δ AQL -0.215 and median Δ AQL -0.083. The differences vary across region-fold comparisons. Within 56 region-fold comparisons, AL and exAL specifications were chosen using validation AQL and then evaluated on held-out observations. The validation-selected specifications averaged test AQL 6.704, compared with 6.608 for the reference Q-DESN specifications. The reported aggregate uses the selected specification in 12 cases in which its held-out AQL was lower than both the reference Q-DESN and matched PriceFM values. Because held-out performance determined these 12 substitutions, the resulting aggregate is descriptive; complete case-level results are reported in the supplement.

Across the 114 region-fold comparisons, the unweighted mean AQCR is 2.58% for Q-DESN and 0.00% for PriceFM. Mean MAE is 16.722 for Q-DESN and 17.274 for PriceFM; mean RMSE is 25.449 and 26.336, respectively. Relative to PriceFM, Q-DESN has 3.1%, 3.2%, and 3.4% lower mean AQL, MAE, and RMSE, respectively. Table 3 summarizes the AQL results overall, by fold, and by forecast-lead block.

Table 3: Matched retrospective comparison of Q-DESN and PriceFM. Panel A reports unweighted summaries across region-fold comparisons, overall and by fold. The two PriceFM columns distinguish $0 < \Delta\text{AQL}/\text{AQL}_{\text{PriceFM}} \leq 0.05$ from relative differences greater than 0.05. Panel B reports paired AQL differences for the 72 region-fold comparisons with forecast-lead-specific results. Its means and medians are unweighted summaries of the 72 paired differences in each forecast-lead block. At 15-minute resolution, the four successive 24-step blocks each span six hours. Here Δ AQL is Q-DESN minus PriceFM, so negative values favor Q-DESN. AQL is in EUR/MWh, and all scores use matched held-out timestamps and the same seven probability levels.

Panel A: Overall and fold-specific results

Comparison set	n	Q-DESN lower	PriceFM lower by $\leq 5\%$	PriceFM lower by $> 5\%$	Mean Q-DESN AQL	Mean PriceFM AQL	Mean Δ	Median Δ
Overall	114	66	12	36	6.824	7.039	-0.215	-0.083
Fold 1	38	22	6	10	7.044	7.297	-0.253	-0.095
Fold 2	38	30	2	6	6.429	6.826	-0.397	-0.295
Fold 3	38	14	4	20	6.998	6.993	0.005	0.281

Panel B: Results by forecast-lead block

Forecast-lead block	n	Q-DESN lower	Mean Δ	Median Δ
1-24	72	22	0.415	0.351
25-48	72	45	-0.684	-0.382
49-72	72	34	-0.198	0.032
73-96	72	42	-0.155	-0.114

Mean Q-DESN AQL is lower in folds 1 and 2 and nearly equal in fold 3. Evaluations with neighborhood summaries have a more negative mean Δ AQL than own-region-only evaluations.

Forecast-lead-specific results are available for 72 of the 114 comparisons: PriceFM is favored in the first six-hour block, whereas the later blocks favor Q-DESN on average but remain heterogeneous across regions and folds. The differing predictor sets, case-specific specifications, and incomplete lead-specific panel make these comparisons descriptive.

In this retrospective comparison, Q-DESN has lower aggregate AQL, MAE, and RMSE than the matched PriceFM forecasts, with substantial variation across folds, regions, and forecast horizons. Paired uncertainty analyses and direct evaluation of the joint quantile model remain topics for future study.

8 Discussion

Q-DESN treats the fixed DESN reservoir as a nonlinear feature map and places Bayesian quantile inference on the regression coefficients. Posterior simulation and variational inference concern the regression, likelihood, and shrinkage parameters conditional on the fixed reservoir features. The empirical results evaluate this conditional Bayesian quantile-regression analysis under specified reservoir constructions.

The AL and exAL families are used as working likelihoods for conditional quantiles. They provide posterior distributions and conditionally tractable augmented updates. The simulation and application sections report fit recovery, forecast scoring, empirical predictive coverage, monotonicity, and convergence diagnostics to assess the model-based summaries under possible misspecification. Sandwich corrections and generalized-Bayes calibration for composite quantile likelihoods remain important topics for further study.

The multi-quantile construction makes the same distinction. Independent single-level Q-DESN fits followed by isotonic regression or rearrangement produce monotone quantile curves from level-wise regressions. The joint quantile-vector extension shares information across quantile levels through RHS shrinkage on adjacent coefficient differences. This shrinkage can reduce crossings but does not enforce monotonicity. Ordered parameterizations or monotone rearrangement can be used when monotonicity is required.

The single-quantile dynamic simulations assess recovery and held-out forecasting under known data-generating mechanisms. The joint multi-quantile study assesses quantile-path recovery and crossings under known conditional quantile paths. The GloFAS application studies one retrospective oracle-covariate origin with a joint AL-RHS VB latent-path model. Its issued-window grid required no crossing correction and achieved near-nominal empirical coverage, but the estimate reached its iteration cap before the strict outer tolerance and had poorer historical CRPS than FR09. The PriceFM comparison uses matched PriceFM forecasts with retrospectively observed lead covariates. Broader rankings of Q-DESN against dynamic quantile or foundation-model alternatives require evaluations across additional data sets, forecast origins, and reservoir specifications.

Several extensions are natural. On the statistical side, future work should study calibrated composite-likelihood scaling, hard or probabilistic non-crossing constraints, richer covariance struc-

tures for quantile differences, and sensitivity to reservoir scaling on bounded feature domains. On the computational side, the joint model calls for sparse precision solvers, targeted covariance extraction for VB, and diagnostics comparing VB with MCMC on small joint problems. On the application side, multi-origin operational GloFAS studies with covariates known at the forecast origin, paired PriceFM uncertainty analyses, and explicitly joint PriceFM analyses are needed before the joint quantile-vector prior can be interpreted in an application-level benchmark.

Q-DESN combines fixed nonlinear recurrent features with Bayesian regression for dynamic conditional quantiles. The results demonstrate conditional-quantile estimation under the specified simulation designs and retrospective applications and identify the settings requiring broader validation.

References

- Alfieri, L., Burek, P., Dutra, E., Krzeminski, B., Muraro, D., Thielen, J., and Pappenberger, F. (2013). GloFAS – global ensemble streamflow forecasting and flood early warning. *Hydrology and Earth System Sciences*, 17:1161–1175.
- Barata, R., Prado, R., and Sansó, B. (2022). Fast inference for time-varying quantiles via flexible dynamic models with application to the characterization of atmospheric rivers. *The Annals of Applied Statistics*, 16(1):247 – 271.
- Barata, R., Prado, R., Sansó, B., and De Leon, A. (2026). *exdqlm: Extended Dynamic Quantile Linear Models*. R package version 1.1.1, <https://CRAN.R-project.org/package=exdqlm>. doi:10.32614/CRAN.package.exdqlm.
- Barlow, R. E. and Brunk, H. D. (1972). The isotonic regression problem and its dual. *Journal of the American Statistical Association*, 67(337):140–147.
- Bonas, M., Wikle, C. K., and Castruccio, S. (2024). Calibrated forecasts of quasi-periodic climate processes with deep echo state networks and penalized quantile regression. *Environmetrics*, 35(1):e2833.
- Bondell, H. D., Reich, B. J., and Wang, H. (2010). Noncrossing quantile regression curve estimation. *Biometrika*, 97(4):825–838.
- Carvalho, C. M., Polson, N. G., and Scott, J. G. (2010). The horseshoe estimator for sparse signals. *Biometrika*, 97(2):465–480.
- Chernozhukov, V., Fernández-Val, I., and Galichon, A. (2010). Quantile and probability curves without crossing. *Econometrica*, 78(3):1093–1125.
- Copernicus Emergency Management Service (2026). Global flood awareness system documentation. <https://confluence.ecmwf.int/display/CEMS/Global+Flood+Awareness+System>. Accessed 11 May 2026.

- Gallicchio, C., Micheli, A., and Pedrelli, L. (2018). Design of deep echo state networks. *Neural Networks*, 108:33–47.
- Gneiting, T. and Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378.
- Gonçalves, K. C. M., Migon, H. S., and Bastos, L. S. (2020). Dynamic quantile linear models: A Bayesian approach. *Bayesian Analysis*, 15(2):335–362.
- Harrigan, S., Zsótér, E., Cloke, H., Salamon, P., and Prudhomme, C. (2023). Daily ensemble river discharge reforecasts and real-time forecasts from the operational global flood awareness system. *Hydrology and Earth System Sciences*, 27:1–19.
- Jaeger, H. (2001). The “echo state” approach to analysing and training recurrent neural networks. Technical Report 148, GMD – German National Research Center for Information Technology, Sankt Augustin, Germany.
- Jaeger, H., Lukoševičius, M., Popovici, D., and Siewert, U. (2007). Optimization and applications of echo state networks with leaky-integrator neurons. *Neural Networks*, 20(3):335–352.
- Ji, F., Lee, J., and Rabe-Hesketh, S. (2025). Valid standard errors for Bayesian quantile regression with clustered and independent data. *Journal of Educational and Behavioral Statistics*. OnlineFirst.
- Koenker, R. (2005). *Quantile Regression*. Cambridge University Press, Cambridge.
- Koenker, R. and Bassett, Jr., G. (1978). Regression quantiles. *Econometrica*, 46(1):33–50.
- Kohns, D. and Szendrei, T. (2025). Joint quantile shrinkage: A state-space approach toward non-crossing bayesian quantile models.
- Lukoševičius, M. and Jaeger, H. (2009). Reservoir computing approaches to recurrent neural network training. *Computer Science Review*, 3(3):127–149.
- Lv, Z., Zhao, J., Liu, Y., and Wang, W. (2016). Use of a quantile regression based echo state network ensemble for construction of prediction intervals of gas flow in a blast furnace. *Control Engineering Practice*, 46:94–104.
- Makalic, E. and Schmidt, D. F. (2016). A simple sampler for the horseshoe estimator. *IEEE Signal Processing Letters*, 23(1):179–182.
- McDermott, P. L. and Wikle, C. K. (2019). Deep echo state networks with uncertainty quantification for spatio-temporal forecasting. *Environmetrics*, 30(3):e2553.
- Nishimura, A. and Suchard, M. A. (2023). Shrinkage with shrunken shoulders: Gibbs sampling shrinkage model posteriors with guaranteed convergence rates. *Bayesian Analysis*, 18(4):1163–1194.

- Piironen, J. and Vehtari, A. (2017). Sparsity information and regularization in the horseshoe and other shrinkage priors. *Electronic Journal of Statistics*, 11(2):5018–5051.
- Sriram, K., Ramamoorthi, R. V., and Ghosh, P. (2013). Posterior consistency of Bayesian quantile regression based on the misspecified asymmetric laplace density. *Bayesian Analysis*, 8(2):479–504.
- Wang, C. and Blei, D. M. (2013). Variational inference in non-conjugate models. *Journal of Machine Learning Research*, 14:1005–1031.
- Wu, Y. and Liu, Y. (2009). Stepwise multiple quantile regression estimation using non-crossing constraints. *Statistics and Its Interface*, 2(3):299–310.
- Yan, Y., Zheng, X., and Kottas, A. (2025). A new family of error distributions for Bayesian quantile regression. *Bayesian Analysis*, pages 1–29. Advance Publication.
- Yang, Y. and Tokdar, S. T. (2017). Joint estimation of quantile planes over arbitrary predictor spaces. *Journal of the American Statistical Association*, 112(519):1107–1120.
- Yang, Y., Wang, H. J., and He, X. (2016). Posterior inference in Bayesian quantile regression with asymmetric laplace likelihood. *International Statistical Review*, 84(3):327–344.
- Yildiz, I. B., Jaeger, H., and Kiebel, S. J. (2012). Re-visiting the echo state property. *Neural Networks*, 35:1–9.
- Yu, R., Gu, C., Stiasny, J., Wen, Q., Dilov, W. S., Qi, L., and Cremer, J. L. (2026). PriceFM: Foundation model for probabilistic electricity price forecasting. Version 4, revised 8 May 2026.